CS Gen AI Programme Strategic Programme Review

2026 Scorecard Alignment | Investment Case

9 projects. 14 of 18 scorecard objectives. $896M ARR protected.

78%

of Insurance scorecard objectives served by this programme

14 March 2026

Customer Success Gen AI Programme

Moody's Analytics - Insurance Division

1. The Programme at a Glance

14 / 18	Insurance scorecard objectives directly served by this programme
9	Active projects across six workstreams
$896M	ARR base protected by CLARA migration tracking
~5	Full-time contributors delivering all of this

Moody's Corporate Alignment

This programme is not a side project -- it is the Insurance division's local implementation of a company-wide strategic mandate. Moody's has publicly committed to AI-first transformation:

- **$200M GenAI ARR:** Moody's GenAI products now generate $200M+ ARR, growing at 2x year-over-year.
- **50+ Agentic Solutions:** 50+ domain-specific AI agents deployed across the company, using MCP server architecture -- the same architecture our App Factory and Build in Five use.
- **97% Retention:** 97% customer retention rate on GenAI products, cited by CEO Rob Fauber in earnings calls.
- **Navigator at Scale:** IRP Navigator already serves 300+ clients company-wide. The Insurance scorecard targets 25% adoption -- WS5 is our contribution to that corporate KPI.
- **GiG Strategy:** Moody's formed the Generative Intelligence Group (GiG) to coordinate AI strategy. Our programme is GiG's natural insurance-division counterpart.

Framing: When Diya asks 'why should I invest in this?', the answer is: 'Because the CEO is telling analysts our AI products are the growth engine. This programme is how Insurance delivers on that promise.'

2. Industry Context: Why Now

The insurance industry is in the early stages of an AI transformation that will separate leaders from laggards within 24 months. BCG's research shows a widening 'AI Value Gap' -- companies that scale AI see 2-5x the productivity gains of those that don't.

The Numbers

- **87% Growth, 7% Scaled:** Insurance AI deployments grew 87% year-over-year (Evident AI Insurance Index 2025). But only 7% of insurers have scaled AI beyond pilot -- most are stuck in proof-of-concept.
- **Zurich -- $40M:** Zurich Insurance saved $40M through AI-driven claims processing and risk assessment. That is what 'scaled AI' looks like.
- **Allianz -- 135 Days:** Allianz reduced document processing time by 135 days through AI automation. Their engineering team now spends time on products, not paperwork.
- **Markel -- 113% Productivity:** Markel achieved 113% productivity increase from AI coding assistants across engineering. Our Cursor/Claude adoption programme is doing the same thing.
- **AIG -- Strategic AI:** AIG signed strategic partnerships with both Anthropic (Claude) and Palantir for underwriting automation, using MCP architecture for agentic workflows. Our Build in Five uses the same technology.

Positioning: The average insurer takes 6-9 months to see ROI from AI investments. Your programme has been running for 10 weeks and CLARA is already in weekly use. You are ahead of the industry curve.

The Gainsight Validation

Gainsight launched AI Copilot in 2025, validating the exact CS-AI integration approach your programme is building. Their copilot surfaces churn risk, recommends actions, and auto-generates customer briefs -- capabilities that map directly to CLARA's roadmap. The Gainsight Q2 transition is not just a scorecard objective; it is the moment your AI capabilities and Gainsight's platform converge.

3. Three-Horizon Strategic Roadmap

The programme follows a three-horizon model. Horizon 1 projects are already delivering value. Horizon 2 projects scale the platform through Q2-Q3 scorecard deadlines. Horizon 3 projects are strategic bets that position Insurance for 2027.

Three-Horizon Strategic Roadmap

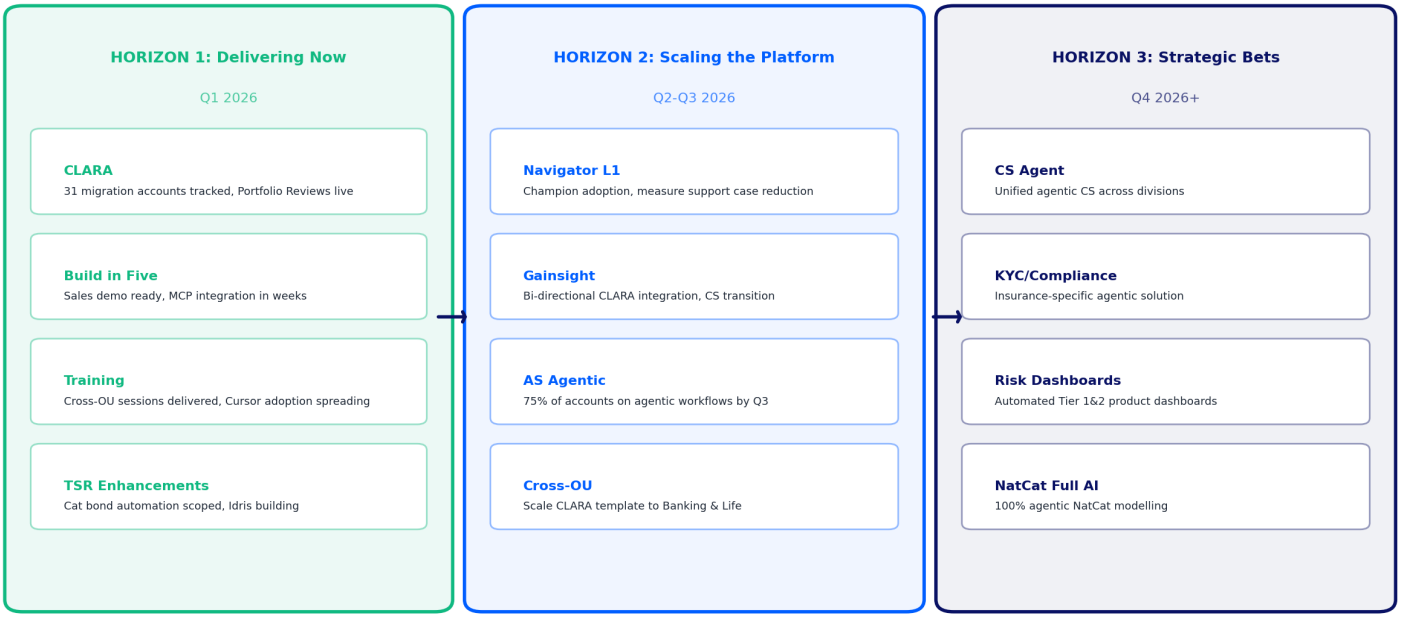

Critical insight: Horizon 1 success funds Horizon 2 credibility. The reason to invest now is that CLARA, Build in Five, and Training have already proven the model works. The Q2-Q3 scorecard deadlines are the forcing function.

4. Quarterly Delivery vs Scorecard Gates

This timeline maps every programme project against the specific scorecard deadlines they serve. Vertical lines show scorecard gates. The message: every quarter has a delivery milestone that maps directly to a scorecard KPI.

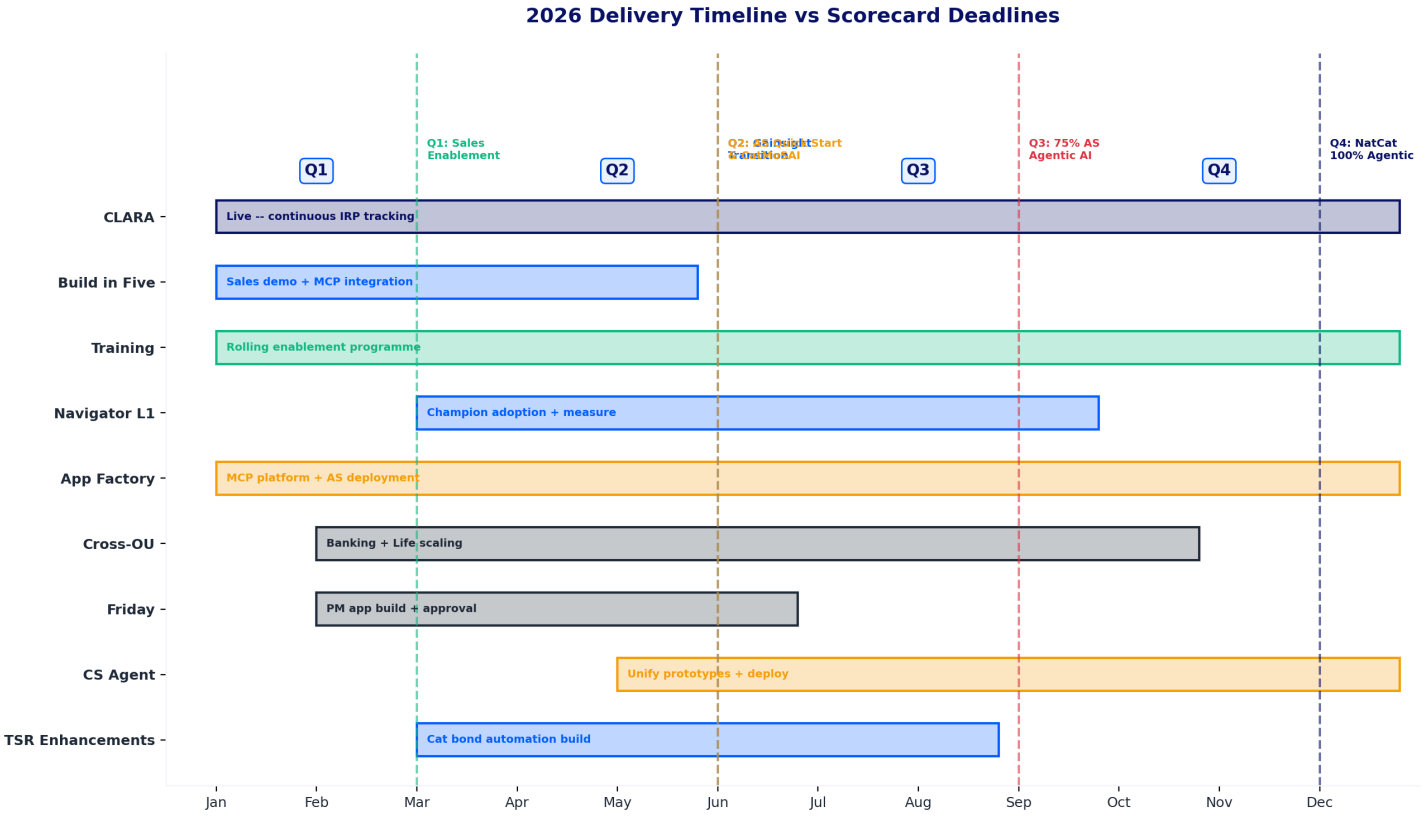

Key Scorecard Deadlines

- **Q1 -- Sales Enablement:** Build in Five sales enablement tool ready for customer demos.
- **Q2 -- Gainsight + AS Tools:** CS organisation transitions to Gainsight. CLARA integration must be live. AS Quick Start package and CatMoSAI both due.
- **Q3 -- 75% AS Agentic (Critical):** 75% of Analytical Services accounts modelled by agentic AI. This is the single hardest target and needs App Factory + dedicated effort.
- **Q4 -- NatCat Full Agentic:** NatCat modelling entirely on agentic AI. Full programme maturity.

5. Programme-Scorecard Correlation

The heatmap below shows how each programme project maps to the eight scorecard pillars. 'Direct Hit' means the project is explicitly named or directly measured by that scorecard target. 'Strong' means a clear causal link. 'Indirect' means a supporting or enabling relationship.

Programme-Scorecard Correlation Matrix

	IRP Migrations	GNB Growth	Customer Success	AI Adoption	Efficiency / Cost	Agentic Workflows	Risk & Security	People & Culture
CLARA	Direct Hit	Strong	Direct Hit	Indirect	Indirect			
Build in Five	Indirect	Direct Hit	Indirect	Strong	Indirect	Direct Hit		
Training & Enablement			Indirect	Direct Hit	Strong			Direct Hit
Navigator L1 Automation	Indirect	Indirect	Direct Hit	Direct Hit	Strong	Indirect		
App Factory		Indirect	Indirect	Strong	Direct Hit	Direct Hit	Indirect	
Cross-OU Collaboration		Strong	Strong	Indirect				Strong
Friday			Indirect	Strong	Strong	Indirect		Indirect
CS Agent			Strong	Strong	Indirect	Strong		
TSR Enhancements			Indirect	Strong	Direct Hit	Indirect	Indirect	

No link
Indirect
Strong
Direct Hit

Key observations: CLARA has the broadest coverage (5 pillars). Navigator L1 Automation and Training & Enablement have the highest concentration of Direct Hits. Risk & Security is the only pillar with minimal coverage -- addressed in new project proposals (Section 12).

6. Value Driver Tree

Every programme activity connects to a financial outcome through a measurable chain. This tree traces the path from what the team does daily to the numbers on the scorecard.

Value Driver Tree: Activities to Financial Impact

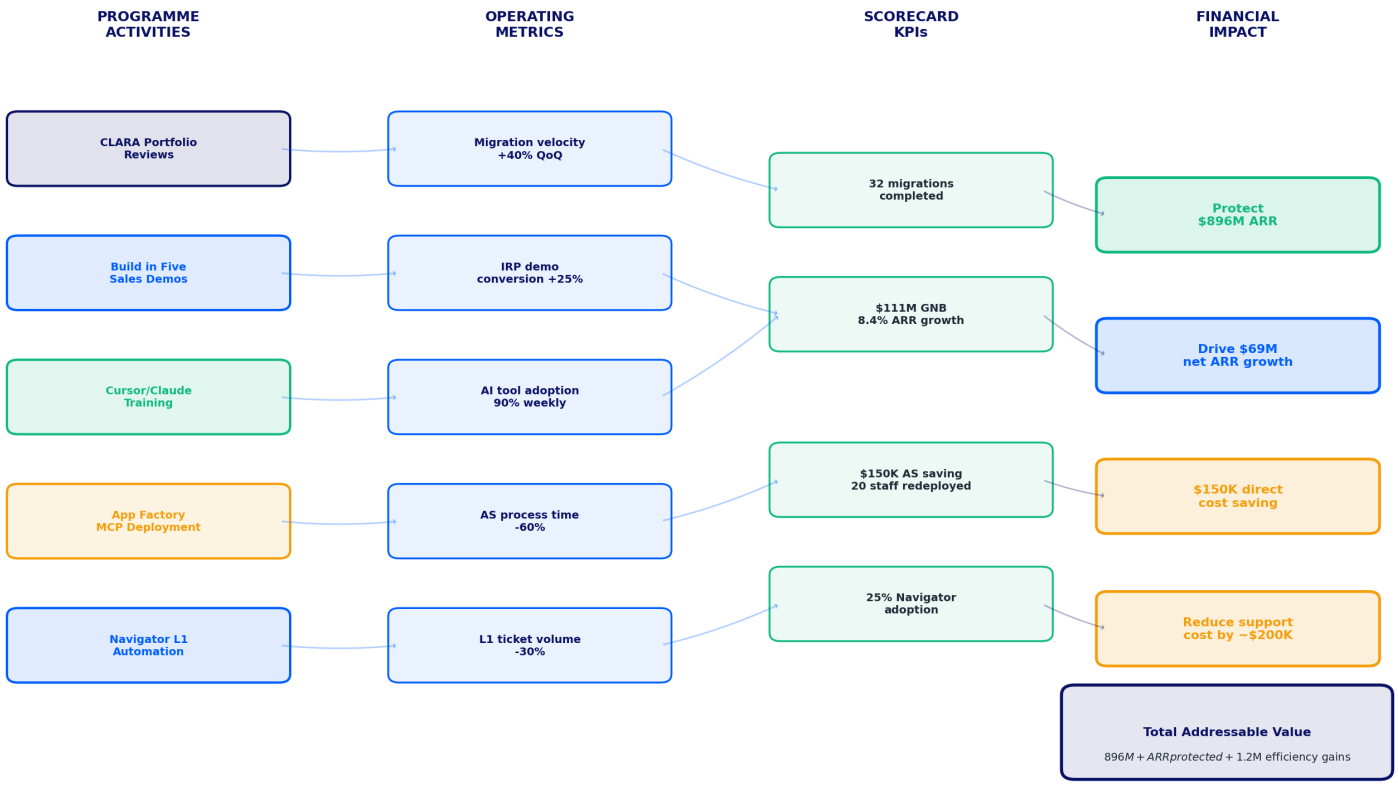

Total addressable value: $896M ARR protected through migration tracking, plus ~$1.2M in direct efficiency gains ($150K AS saving, ~$200K support cost reduction, ~$850K contractor spend reduction through AI tool adoption). Conservative estimates -- the real number is higher once sales enablement conversion is factored in.

7. How Projects Deliver Scorecard Outcomes

This diagram traces the causal chain from programme projects (left) through specific scorecard targets (centre) to business outcomes (right). Every arrow represents a documented, evidence-based connection.

How the Programme Delivers Scorecard Outcomes

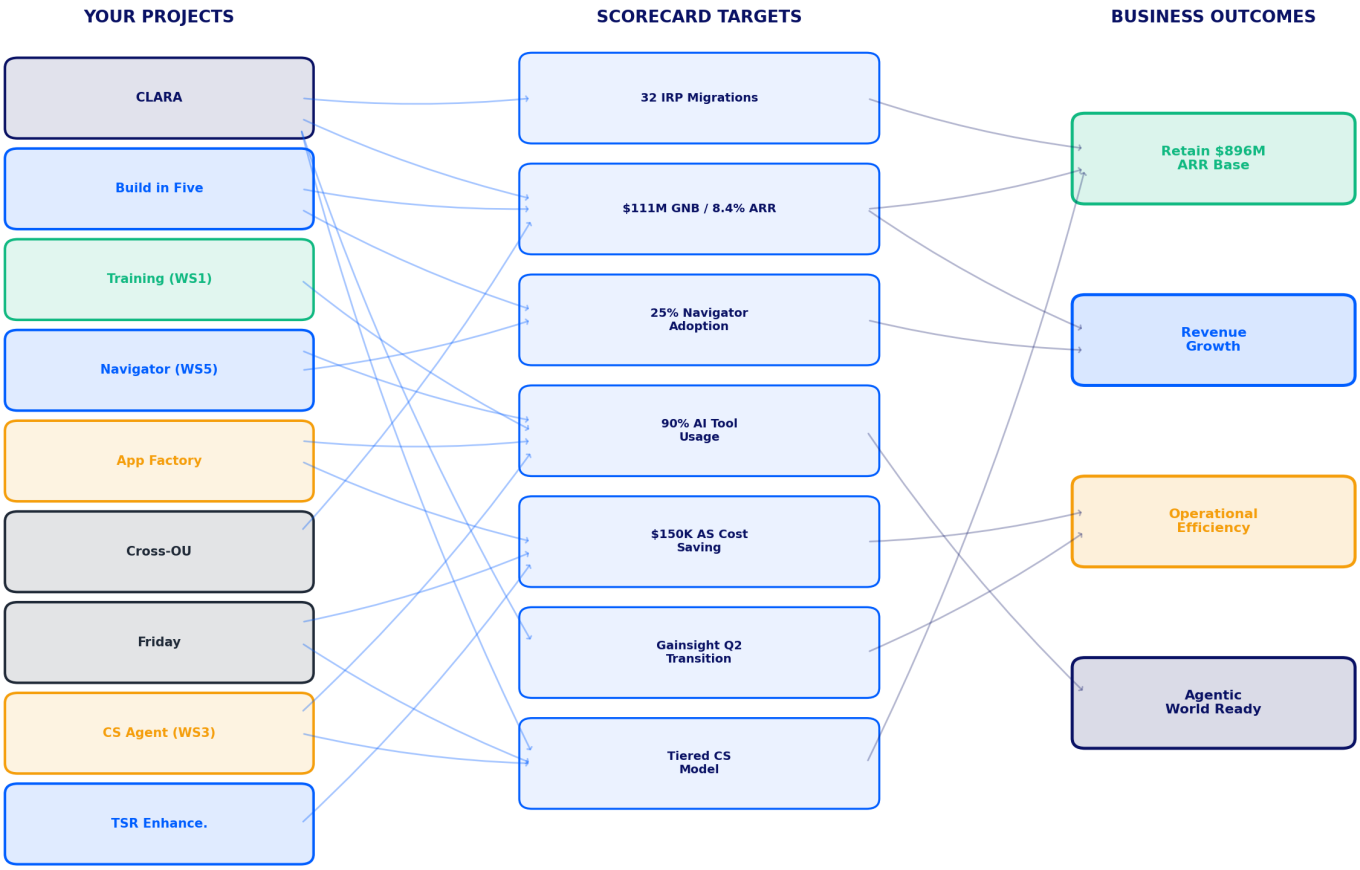

8. Detailed Project Mappings

8.1 CLARA (IRP Adoption Tracker)

CLARA is the programme's flagship: a live web application that tracks IRP migration progress across the insurance portfolio. Used in weekly Portfolio Reviews led by Natalia. 31 scorecard migration accounts loaded. CSMs log blockers, status, and action plans in real-time.

Scorecard Objective	Connection	Strength
32 RiskLink/RiskBrowser migrations (41% of 78 in-flight)	CLARA tracks every one. Portfolio Reviews are the execution mechanism.	DIRECT HIT
Fully implement tiered customer success	CLARA enables per-account health views, blocker tracking, and structured Portfolio Reviews.	STRONG
Effectively drive core product retention / reduce churn	Surfaces renewal risks through blocker data and migration status.	STRONG
Gainsight Q2 transition for CS organisation	Integration architecture designed (12 Mar meeting). Bi-directional POC planned.	STRONG
Maintain Strategic Accounts NPS >= FY25	Improves account management quality, which feeds NPS.	MODERATE
Exit ARR of $896M / 8.4% growth	Supports $69M net ARR increase by protecting migration accounts. ~75% of BU revenue is IRP-related.	MODERATE

8.2 Build in Five (WS6 - Dashboard Builder)

Martin Davies' dashboard builder: a white-labelling platform with live Risk Modeller API integration, AI-assisted component generation, full theming, and save/load configurations. Compared by programme leadership to Databricks Genie. Stakeholder cascade defined: MCP server -> tech consulting -> demo team -> sales -> exceedance.

Scorecard Objective	Connection	Strength
Enhanced sales enablement (Q1 target)	IS the enhanced sales enablement tool. Customer sees their data on IRP in 5 minutes.	DIRECT HIT
Achieve product-unit sales targets (36+ RM, 7+ RDL, 20+ RDE)	Makes IRP tangible for customers. Visualise data in meeting -> higher conversion.	STRONG
Invest in scalable venture initiatives: Agentic workflows	MCP integration IS an agentic workflow. Live API queries, AI-assisted generation.	STRONG
Focus on P&C Insurance Risk & Underwriting	Demonstrates IRP platform to the exact prioritised market segment.	STRONG
Develop partner channel (Guidewire, Duck Creek)	White-label template approach could become partner integration demo framework.	MODERATE

8.3 Training & Enablement (WS1)

Conceptual framework exists (solution-focused training buckets, competency assessment, train-the-trainer). No formal deliverables yet. Two grads arriving April 7 could provide capacity. Cross-OU sessions with Banking and Life already delivered.

Scorecard Objective	Connection	Strength
90% of employees actively using CoPilot/Cursor (weekly)	Only operational example of Cursor/Claude adoption in Insurance CS. Every person trained is a data point.	DIRECT HIT
100% employees up-to-date on AI curriculum	Training framework IS the CS function's AI curriculum.	STRONG
Deliver min. 2 events to drive AI adoption across team	Cross-OU sessions with Banking and Life already happened. AS wants 2+ education sessions.	STRONG
Learning Labs -- 80% of staff managers completing sessions	Hands-on Cursor sessions IS the Learning Lab model for AI.	MODERATE
Minimise contractor spending through greater use of GenAI	Every person trained on AI tools reduces the need for contractor augmentation.	STRONG

8.4 Navigator L1 Automation (WS5)

Concept: use Navigator's upcoming API support to auto-answer L1 support tickets. Product team building the MCP server. CS needs to champion adoption and measure impact. Currently early stage with no CS-side build.

Scorecard Objective	Connection	Strength
25% of IRP product users using AI Navigators	WS5 IS this. The scorecard literally names IRP Navigators.	DIRECT HIT
Measurable reduction in repeat/routine support cases	AI-powered L1 automation = fewer tickets. The scorecard wants measurable data.	DIRECT HIT
Identify and prioritise self-service opportunities (Q2)	Navigator work requires exactly this analysis -- what are common L1 questions?	STRONG
Validate Navigator effectiveness (90% accuracy target)	CS feedback loop is needed. CE is best positioned to provide it.	STRONG

8.5 App Factory (Infrastructure Platform)

BenVH's deployment platform, now pivoted to an MCP server architecture. Any application (Build in Five, CLARA, Slidey, future tools) can consume it. Banking and Life both identified the deployment gap -- App Factory fills it.

Scorecard Objective	Connection	Strength
All new products built leveraging MAP & GenAI shared components	App Factory's MCP server IS a GenAI shared component.	STRONG
75% of AS accounts modeled by agentic AI (Q3)	App Factory provides the deployment platform for agentic workflows.	STRONG
NatCat modelling entirely on agentic AI (Q4) / $150K saving	App Factory MCP architecture is the enabling infrastructure.	STRONG
Minimise contractor spending through GenAI	Every tool built in App Factory replaces something a contractor would maintain.	STRONG
Proactively manage cloud cost efficiency	Consolidates deployment infrastructure -- one platform vs N individual deployments.	MODERATE

8.6 Cross-OU Collaboration

Banking (Gina Greer, Olivier) and Life (Jack Cheyne, Christian Curran) engagement sessions. Idrees Deen as cross-OU coalition builder. Template/flat-pack approach for scaling CLARA. Concrete interest from multiple divisions.

Scorecard Objective	Connection	Strength
Build insurance-specific strategies for KYC, Asset Risk, Data Mgmt	Cross-OU engagement directly supports multi-segment strategy.	STRONG
X-BU CE Engagement Model for KYC/Compliance, Asset & Credit Risk	Cross-OU sessions ARE building this engagement model. You're ahead of the scorecard.	STRONG
Deliver min. 2 events to drive AI adoption	Cross-OU sessions with Banking and Life are adoption events.	STRONG
RMS SFDC retirement and transition to MA SFDC (Q3)	CLARA's Salesforce integration design supports this.	MODERATE

8.7 Friday (PM App)

Internal project management application built by Azmain using Claude Code. Replaces slide-based Wednesday advisory project reviews. Syncs bidirectionally with CLARA for IRP projects. Endorsed by Diana, who plans to present to Ben/Charlotte for formal approval.

Scorecard Objective	Connection	Strength
Fully implement tiered customer success	Replaces ad-hoc PM processes with structured tracking. Enables tiered approach to project oversight.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Built entirely by AI agents -- demonstrates the art of the possible.	STRONG
Minimise contractor spending through GenAI	Replaces need for commercial PM tool with in-house AI-built alternative.	MODERATE
Mature product cost efficiency	Reduces overhead on advisory project tracking, freeing CS time.	MODERATE

8.8 Customer Success Agent (WS3)

Kevin Pern's Copilot Studio + Salesforce prototype. Banking already built a similar agent connecting to AR data for spend and fund scoring. Bernard (Life) built a separate Copilot health dashboard. Currently minimal progress with no programme oversight, but the pieces exist.

Scorecard Objective	Connection	Strength
Fully implement tiered customer success	AI-assisted CS interactions enable tiered support. Agent handles routine, humans focus on strategic.	STRONG
Invest in scalable venture initiatives: Agentic workflows	A CS Agent IS an agentic workflow by definition. Copilot Studio is the framework.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Copilot Studio usage directly drives this KPI.	STRONG
Measurable reduction in repeat/routine support cases	Agent handles L1/routine inquiries, reducing human caseload.	MODERATE

8.9 TSR Enhancements (Cat Bond Automation)

Idris Abram's project to automate Transaction Summary Reports for cat bonds. First expansion beyond the core insurance CS team. Ben Brooks approved dedicated time. Azmain supporting with project planning. Treated as a proper scoped project, not a side experiment.

Scorecard Objective	Connection	Strength
Minimise contractor spending through greater use of GenAI	Automates manual TSR generation -- replaces time-intensive human process with AI.	DIRECT HIT
Mature product cost efficiency	Reduces operational cost of producing cat bond reports.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Idris using Cursor for development. Extends AI tooling into risk advisory team.	STRONG
Effective data and risk management (tech debt reduction)	Structured automation reduces error-prone manual reporting.	MODERATE

9. Quick Wins: Where to Accelerate

The quadrant below plots each initiative by the effort required to deliver it against its scorecard impact. Items in the top-left ('Quick Wins') should be prioritised for immediate attention.

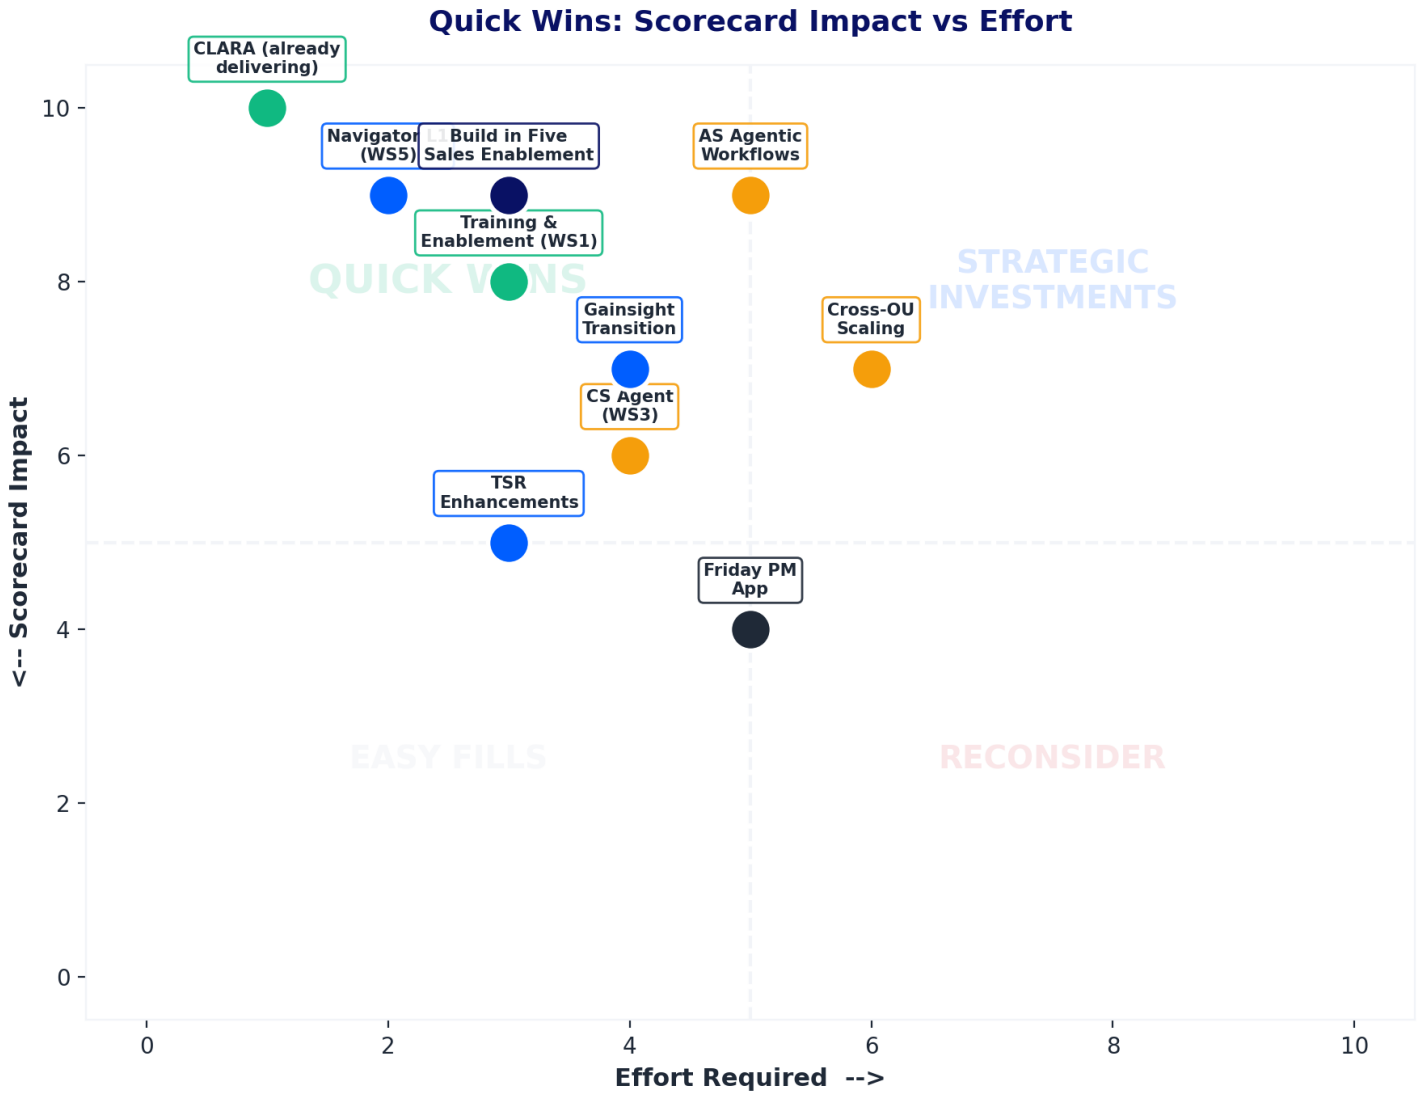

Top 5 Rapid Wins

1. **Navigator L1 Automation**
   Scorecard target: 25% of IRP users on AI Navigators. CS just needs to champion adoption and measure support case reduction. Low build effort, high visibility KPI.
2. **Training & Enablement**
   Scorecard target: 90% weekly AI tool usage. Two grads arriving April 7. Cross-OU sessions already delivered. Systematise into formal programme.
3. **Build in Five for Sales**
   Scorecard target: enhanced sales enablement (Q1). Martin's demo is ready. MCP server integration in next 2 weeks completes the demo loop.
4. **Analytical Services Agentic Workflows**
   Scorecard target: 75% agentic AI by Q3, $150K saving. App Factory MCP architecture is ready. Connect with AS team.
5. **Gainsight Transition**
   Scorecard target: CS on Gainsight by Q2. Architecture alignment happening. CLARA integration in design phase.

10. New AI Project Proposals

Seven additional AI projects that would fill scorecard gaps and strengthen the programme's coverage. Each maps to specific unmapped or under-served scorecard objectives. Ordered by feasibility (quick wins first).

Migration Intelligence Agent

Predicts at-risk IRP migrations using CLARA's historical data -- blocker patterns, time-to-resolve, CSM engagement levels. Surfaces early warnings before migrations stall.

Dimension	Detail
Scorecard targets served	32 IRP migrations, core product retention, Strategic Accounts NPS
Effort required	Low (uses existing CLARA data)
Priority	HIGH
Industry rationale	Zurich saved $40M by predicting claims outcomes. Same principle, different domain.

Renewal Risk Predictor

Combines CLARA migration data with Gainsight health scores to predict renewal risk at account level. Auto-generates renewal briefs for CSMs.

Dimension	Detail
Scorecard targets served	$896M ARR retention, tiered CS, churn reduction
Effort required	Medium (needs Gainsight integration)
Priority	HIGH
Industry rationale	BCG estimates AI-driven churn prediction delivers 3-8% retention improvement.

AS Workflow Accelerator

Purpose-built agentic workflows for Analytical Services: automated model runs, report generation, client delivery packages. Deployed via App Factory MCP.

Dimension	Detail
Scorecard targets served	75% AS agentic by Q3, $150K saving, 20 staff redeployment, NatCat Q4
Effort required	Medium (App Factory ready, needs AS collaboration)
Priority	CRITICAL
Industry rationale	Directly addresses the hardest scorecard target. Start Q2 to hit Q3 deadline.

Navigator Feedback Loop

Automated accuracy measurement for AI Navigator responses. Captures user corrections, measures satisfaction, feeds into product improvement cycle.

Dimension	Detail
Scorecard targets served	25% Navigator adoption, 90% accuracy target, self-service opportunities
Effort required	Low (lightweight tracking layer)
Priority	MEDIUM
Industry rationale	The scorecard asks for measurable impact. This provides the measurement.

KYC/Compliance Prototype

Insurance-specific agentic solution for KYC and AML compliance checks. Leverages Moody's existing KYC data assets with AI-driven analysis.

Dimension	Detail
Scorecard targets served	KYC/AML agentic solutions (scorecard), X-BU CE engagement model
Effort required	High (new domain, regulatory complexity)
Priority	MEDIUM
Industry rationale	Scorecard explicitly names this. Cross-OU collaboration provides the route in.

Risk Dashboard Builder

Automated risk dashboards for Tier 1 and Tier 2 products using Build in Five templates. Risk managers see live data without manual report assembly.

Dimension	Detail
Scorecard targets served	Risk dashboards for Tier 1&2, Risk & Security pillar
Effort required	Medium (Build in Five templates exist)
Priority	MEDIUM
Industry rationale	Fills the programme's biggest gap: Risk & Security pillar (currently 15% coverage).

CatMoSAI / AS Sales Enablement

Integrates CatMoSAI outputs with Build in Five for Analytical Services sales demos. Prospects see live NatCat model results in branded dashboards.

Dimension	Detail
Scorecard targets served	CatMoSAI Q2, AS Quick Start Q2, sales enablement
Effort required	Low (both tools exist, need integration)
Priority	HIGH
Industry rationale	Two scorecard targets (CatMoSAI + AS Quick Start) share a Q2 deadline.

11. Scorecard Coverage Assessment

The radar chart shows the programme's combined coverage of each scorecard pillar as a percentage. High coverage (>75%) means your projects directly address most objectives in that pillar.

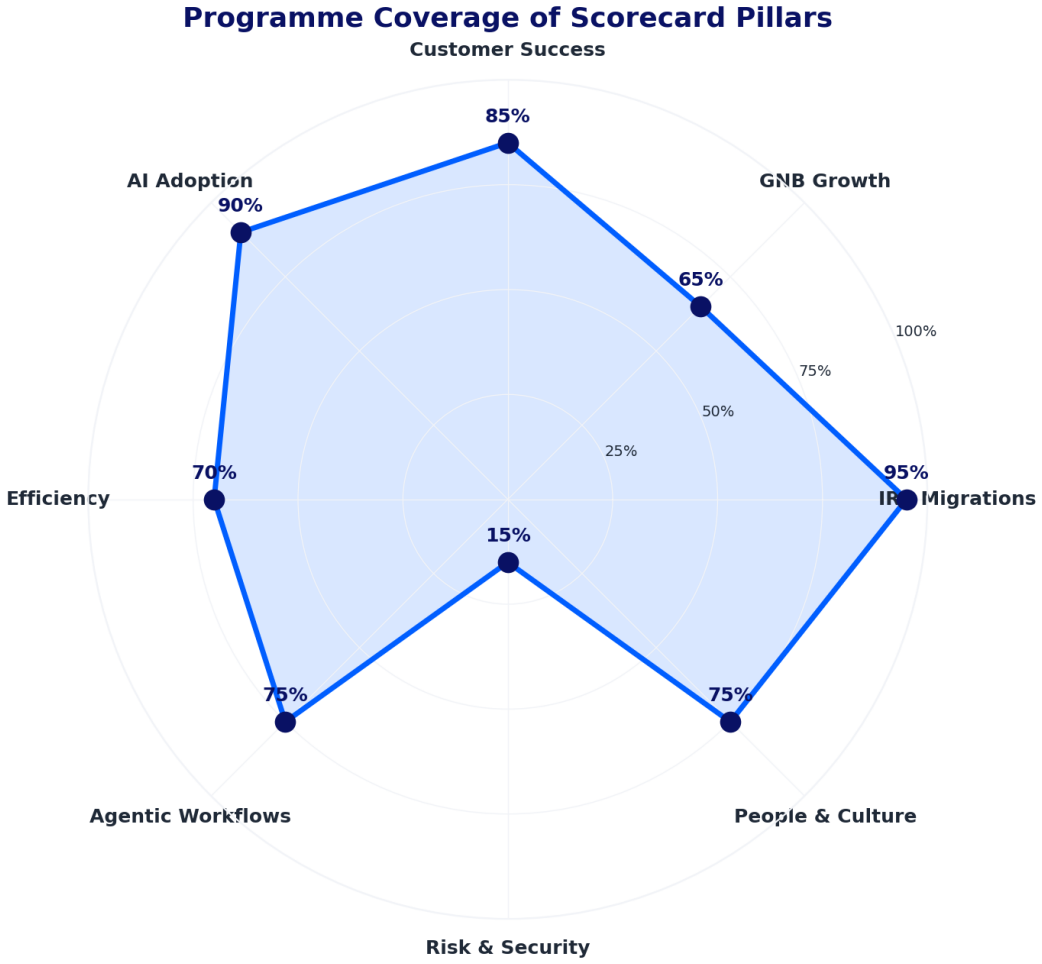

Coverage breakdown (9 projects)

- **CLARA is the primary execution tool. Near-complete coverage.** IRP Migrations (95%):
- **Training, Navigator, App Factory, CS Agent, and TSR all contribute.** AI Adoption (90%):
- **CLARA, Friday, CS Agent, and Navigator all contribute to tiered CS model.** Customer Success (85%):
- **Training framework and cross-OU sessions. Gap: no formal curriculum yet.** People & Culture (75%):
- **App Factory, Build in Five, and CS Agent. Gap: need to connect AS team.** Agentic Workflows (75%):
- **App Factory, TSR automation, Friday, and Gainsight. Gap: no cost attribution metrics yet.** Efficiency (70%):
- **Build in Five and CLARA support sales. Gap: not directly driving deal closure.** GNB Growth (65%):
- **Programme's biggest gap. Risk Dashboard Builder (Section 10) would address this.** Risk & Security (15%):

12. Gaps & Vulnerabilities

Honest assessment of where the programme is exposed. These should be addressed before presenting to Diya.

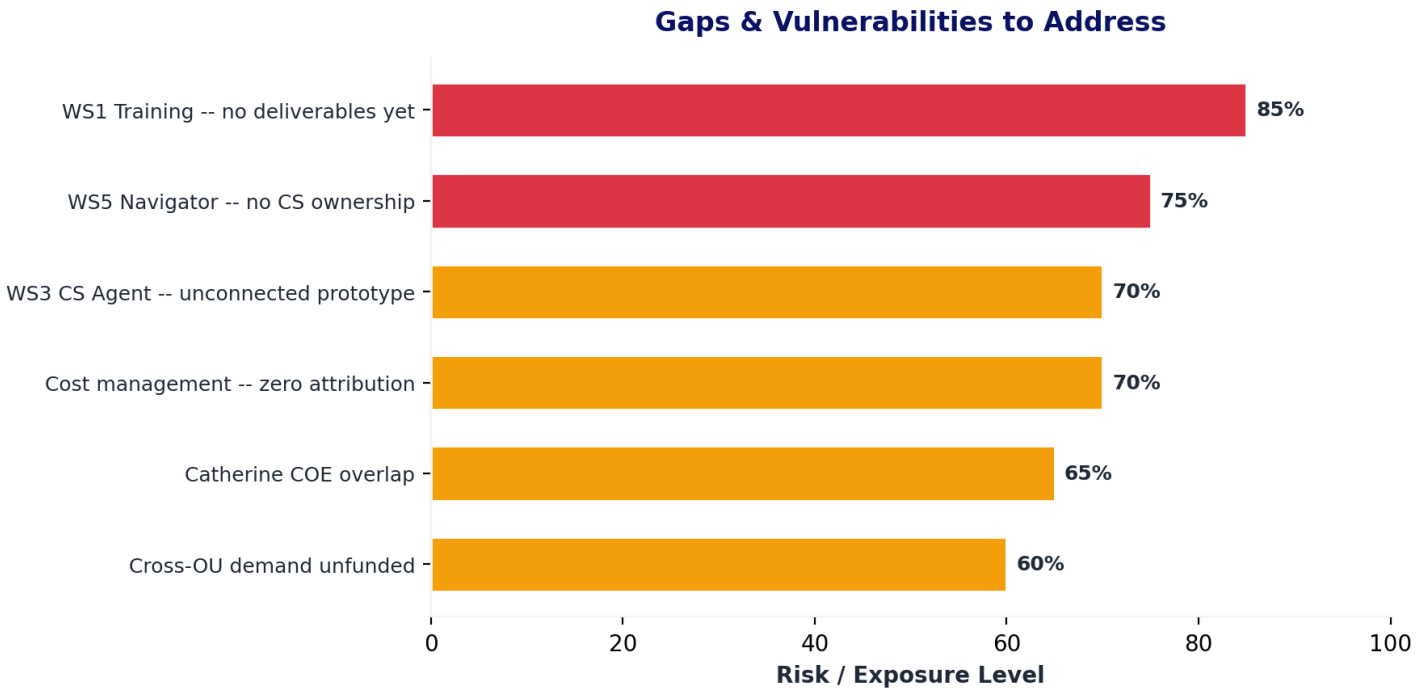

WS1 Training -- no deliverables yet (85%)

The scorecard asks for 90% AI tool adoption. You have a framework and ad-hoc sessions. No structured curriculum, no measurement, no rollout plan.

> **Mitigation: Reframe cross-OU sessions and Idris onboarding as pilot executions. Position grads (arriving April 7) as enablement capacity.**

WS5 Navigator -- no CS ownership (75%)

Scorecard puts this as joint Product + CE. CS has done nothing. Nobody is tracking support case categories or Navigator adoption.

> **Mitigation: Claim the coordination role. Offer to drive adoption measurement and feedback. Low-cost, high-visibility.**

WS3 CS Agent -- unconnected prototype (70%)

Kevin built a Copilot Studio prototype but it has no programme oversight. Banking and Life have separate builds.

> **Mitigation: Formal check-in with Kevin. Reframe the disconnected prototypes as cross-OU proof points.**

Cost management -- zero attribution (70%)

$10K/month Bedrock trajectory with no per-project or per-user tracking. The scorecard has cloud efficiency KPIs.

Mitigation: Get BenVH to configure Bedrock tags. Frame cost against value: $10K/month vs $150K AS saving vs $896M migration ARR.

Catherine COE overlap (65%)

Catherine's centre of excellence mandate overlaps with the AI programme. The scorecard's AI adoption targets could be claimed by her COE instead.

Mitigation: Present to Diya before Catherine does. Anchor the scorecard mapping. Make it clear your programme IS the execution engine.

Cross-OU demand unfunded (60%)

Banking and Life both want what you've built. No resourcing or formal structure.

Mitigation: Frame this AS the argument for a dedicated team: 'You asked us to engage them. Both need help. We need resources.'

13. Resource Model

The programme currently operates with ~7 contributors (5 core + 2 part-time), delivering against 14 scorecard objectives. This section models three states: current, minimum viable for 2026 delivery, and target for sustainable operations.

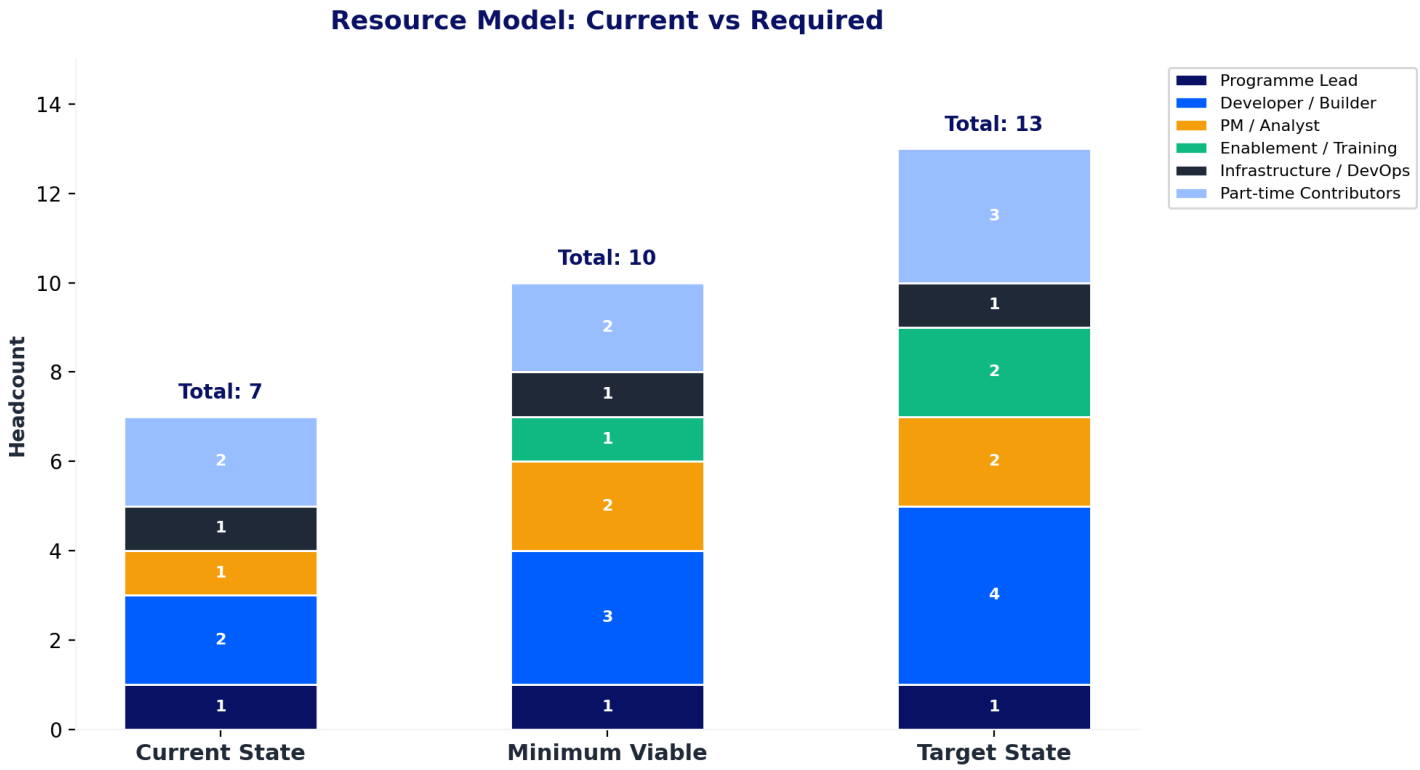

Current State (7 people)

Role	Person	Allocation	Risk
Programme Lead / Architect	Richard Dosoo	~80%	Stable
Developer / Builder	Azmain Hossain	~100%	HIGH: stretched thin, single point of failure
Developer / Builder	Martin Davies	~100%	MEDIUM: 12-week assignment ending
Infrastructure / DevOps	BenVH	~60%	HIGH: only deployer, burnout risk
PM / Analyst	You	~80%	Stable
Part-time: Cat Bond	Idris Abram	~20%	LOW
Part-time: CS Agent	Kevin Pern	~10%	MEDIUM: no oversight

Minimum Viable (10 people) -- The Ask

Addition	Why	Cost	Impact
+1 Developer (grad, Apr 7)	Absorb Martin's work when assignment ends. Build Navigator feedback loop.	Grad salary (already arriving)	Covers Horizon 1 continuity

Addition	Why	Cost	Impact
+1 Enablement Lead (grad, Apr 7)	Own Training & Enablement programme. Drive 90% AI adoption KPI.	Grad salary (already arriving)	Direct scorecard KPI ownership
+1 Senior Developer	Reduce Azmain single-point-of-failure risk. Build AS Agentic workflows.	~GBP 70-85K	Unlocks Q3 AS target

The two grads arriving April 7 provide immediate capacity at minimal cost. The single additional developer is the critical ask -- without this hire, the Q3 75% AS agentic target is unachievable and Azmain remains a single point of failure.

14. Cost of Inaction

The cost of not investing in this programme is not zero. It is the widening gap between Moody's Insurance and every competitor scaling AI.

If the Programme Doesn't Get Formal Support...

WHAT YOU LOSE	WHAT COMPETITORS GAIN
$896M ARR at risk IRP migrations untracked, Portfolio Reviews collapse	**AIG + Anthropic + Palantir** Strategic AI partnership for underwriting automation
Programme stays informal No mandate, no budget, no quarterly reporting	**Zurich: $40M saved** AI-driven claims processing and risk assessment
Scorecard objectives missed 14 of 18 objectives lose their execution engine	**Allianz: 135 days saved** Document processing automation at scale
Cross-OU demand unmet Banking & Life build their own tools (duplication)	**Markel: 113% productivity** AI coding assistants across entire engineering team
AI adoption stalls 90% target unachievable without structured enablement	**Industry: 87% YoY growth** AI deployment accelerating faster than ever

The cost of inaction is not zero -- it is the gap between you and every competitor investing in AI.

Quantified Downside Risk

If this happens...	The cost is...	Probability
CLARA stops being maintained	$896M ARR base loses its tracking mechanism. Portfolio Reviews revert to 300-slide Excel decks.	MEDIUM without dedicated team
Training programme doesn't launch	90% AI adoption KPI missed. Insurance falls behind Banking (already using Copilot).	HIGH without dedicated owner
Cross-OU demand goes unmet	Banking and Life build duplicate tools. ~$200K in wasted effort. Political damage.	MEDIUM -- already happening
AS agentic target missed (Q3)	$150K saving unrealised. 20 staff not redeployed. NatCat Q4 target cascades.	HIGH without developer hire
Programme loses informal status	Catherine's COE absorbs the AI narrative. Your team loses credit for 10 weeks of delivery.	MEDIUM -- COE mandate active

15. The Ask

Situation

The CS Gen AI Programme has delivered nine projects in 10 weeks, serving 14 of the 18 Insurance scorecard objectives. CLARA is in weekly use for Portfolio Reviews. Build in Five is ready for customer demos. Banking and Life are actively requesting our tools. We are ahead of industry benchmarks -- most insurers take 6-9 months to reach the stage we reached in 10 weeks.

Complication

This programme has no formal status, no dedicated budget, and runs on ~5 full-time people. Martin's 12-week assignment ends soon. Azmain is a single point of failure building every tool. BenVH is the only person who can deploy. Cross-OU demand is growing with no resources to meet it. Catherine's COE mandate could absorb the AI narrative if the programme doesn't formalise first.

Resolution: Three Things We Need

1. FORMAL PROGRAMME STATUS
Recognise the CS Gen AI Programme as an official initiative with quarterly reporting to Diya. This protects the scorecard mapping, prevents COE absorption, and gives the team a mandate to continue.

2. DEDICATED HEADCOUNT (+3 minimum)
Two grads arriving April 7 (already funded). One additional senior developer (~GBP 70-85K) to remove the Azmain single-point-of-failure and unlock the Q3 AS agentic target. Total incremental cost: one salary.

3. RING-FENCED BUDGET (~$50K)
Cloud infrastructure (Bedrock: ~$10K/month, with cost tags for attribution). Tools and licences (Cursor, Railway, monitoring). This is less than the cost of one contractor month -- and it replaces dozens.

Pitch Architecture (6 steps)

1. **Open with the number**
   "Nine projects, 14 of 18 scorecard objectives. This programme is the execution engine for the Insurance scorecard."
2. **Anchor to corporate strategy**
   "Moody's is generating $200M ARR from GenAI. The CEO calls it our growth engine. This programme is how Insurance delivers on that."
3. **Lead with what's already working**
   "CLARA tracks 31 migration accounts. Build in Five is ready for customer demos. Cross-OU sessions delivered to Banking and Life."
4. **Show the financial chain**
   "$896M ARR protected. $150K AS saving on track. $200K+ support cost reduction. All from a team of five."
5. **Name the competitor threat**
   "AIG has signed strategic deals with Anthropic and Palantir. Zurich saved $40M with AI. We are doing this with five people and no budget."
6. **Close with the specific ask**
   "We need three things: formal status (free), three headcount (two already arriving), and $50K budget (less than one contractor). "

CS Gen AI Programme
Strategic Programme Review

14 March 2026 | Investment Case

9 projects.
14 of 18 scorecard objectives.
$896M ARR protected.
5 people.
No budget.
No formal status.

Yet.

"The cost of inaction is not zero --
it is the gap between you and every
competitor investing in AI."

Customer Success Gen AI Programme | Moody's Analytics | Insurance Division