CS Gen AI Programme Strategic Programme Review

2026 Scorecard Alignment | Investment Case

9 projects. 14 of 18 scorecard objectives. $896M ARR protected.

78%

of Insurance scorecard objectives served by this programme

14 March 2026

Customer Success Gen AI Programme

Moody's Analytics - Insurance Division

1. The Programme at a Glance

14 / 18	Insurance scorecard objectives directly served by this programme
9	Active projects across six workstreams
$896M	ARR base protected by CLARA migration tracking
~5	Full-time contributors delivering all of this

Moody's Corporate Alignment

This programme is not a side project -- it is the Insurance division's local implementation of a company-wide strategic mandate. Understanding where Moody's stands on AI -- and how aggressively it is moving -- is essential context for anyone deciding whether to invest in this programme.

1a. Moody's Corporate AI Strategy: Deep Dive

The CEO Mandate: AI as the Growth Engine

CEO Rob Fauber has made AI the centrepiece of Moody's growth story. In early 2023, he launched a company-wide GenAI transformation by showing the board a deepfake video of himself delivering a fictitious earnings call -- demonstrating AI capabilities so viscerally that the board approved an aggressive AI strategy immediately. He describes the approach as 'sprinting into fog' -- accepting uncertainty while pursuing substantial rewards. The numbers validate the sprint:

- **$7.7B Revenue, $200M GenAI ARR:** Moody's total 2025 revenue: $7.7 billion (9% YoY growth, record year). Moody's Analytics contributed $3.6 billion. GenAI products generate $200M+ ARR, growing at 2x the rate of the overall business. 40% of MA products by ARR now include GenAI.
- **Research Assistant (Fastest Ever):** Moody's Research Assistant -- launched December 2023 -- became the most rapidly adopted product in Moody's 115-year history. It saw a 300% increase in recurring users in its second quarter. Users access 60% more data while completing tasks 30% faster. Over 90% of interactions are high-value analytical tasks.
- **97% Retention:** 97% customer retention rate on GenAI products (Q4 2025 earnings call). This cohort grows at roughly 2x the rate of the rest of the customer base. Customers who use AI features do not leave -- they pay more.
- **67% Uplift, 35-40% Automation:** Two-thirds of CreditLens renewals are converting to the AI-enabled Lending Suite, generating a 67% revenue uplift per customer. A Tier 1 US bank reported 35-40% of credit memos now auto-generated, saving 'hundreds and thousands of hours.'
- **53,000% Query Growth:** Customer database queries have grown 53,000% over the last 13 years. Moody's now receives nearly a billion hits per day. Fauber: 'We oftentimes talk about black swan events. But now it feels like we have a variety of black swans happening all the time. I tend to think of this as a new era.'

Why this matters: The CEO is telling Wall Street that AI is the growth engine. Every division is expected to deliver AI-powered products and efficiency gains. The Insurance scorecard exists because of this mandate. This programme IS how Insurance delivers on the CEO's promise.

The HBR Case Study: 'How a Legacy Financial Institution Went All In on Gen AI'

In March 2025, Harvard Business Review published a case study by Professor Toby Stuart (UC Berkeley Haas) documenting how Moody's went 'all in' on generative AI. This is not a puff piece -- it is an academic case study that reveals the internal playbook Fauber used to transform a 115-year-old institution:

- **Three Principles:** Fauber set three guiding principles: (1) Make Everyone an Innovator -- AI access democratised to all 14,000 employees, not just tech teams; (2) Build on New Ideas -- a 'yes, and...' mindset that includes legal and compliance from day one; (3) Deliver Impact -- only projects aligned with measurable business outcomes proceed.
- **14,000 Innovators:** Moody's deployed an internal AI tool (CoPilot) to all 14,000 employees. It went from zero to 4,000 users in under 30 days. Employees discovered applications across financial analysis, document processing, sales pitches, and data aggregation -- generating millions of AI prompts and hundreds of use cases.
- **95% Training Target:** Fauber targeted 95% completion of a bespoke AI training programme, with bonus incentives for participation. He held town halls to transform scepticism into enthusiasm. This is the model our WS1 (Training & Enablement) should follow.

- **$750M Disruption Risk:** HBR noted that Moody's acknowledged the risk: its traditional report-based revenue model generates $750M annually and could be disrupted by AI. Rather than protecting the old model, Fauber chose to cannibalise it first. Our programme faces the same choice with legacy IRP workflows.

The HBR playbook maps directly to our programme: (1) democratise AI access (WS1 Training), (2) embed AI in existing workflows (CLARA, Navigator), (3) measure impact (scorecard alignment), (4) move fast (Build in Five). We are not inventing a new approach -- we are executing the CEO's proven playbook at the Insurance division level.

CreditLens: The Proof Point for AI-Enhanced Products

CreditLens is Moody's lending platform and the most concrete example of how AI transforms product economics. Understanding what happened with CreditLens shows what is possible for CLARA and the Insurance product suite:

- **20% Revenue Growth:** CreditLens is approaching 20% revenue growth, driven almost entirely by AI feature adoption. Customers are paying more because AI makes the product materially more valuable.
- **35-40% Auto-Generation:** A Tier 1 US bank reported that CreditLens AI now generates 35-40% of their credit memos automatically, saving what they described as 'hundreds and thousands of hours' of analyst time.
- **80% Faster Decisions:** Banks in APAC and the Middle East using CreditLens AI report 80% reduction in decision time and 15x improvement in loan processing throughput. These are not incremental improvements -- they are step-change transformations.
- **Pricing Power:** The AI features are so effective that two-thirds of renewals convert to the AI-enabled tier, generating 67% more revenue per seat.

Why this matters to the programme: CreditLens proves the business model. When Moody's embeds AI into an existing product, customers pay more, stay longer, and get dramatically better outcomes. CLARA is doing the same thing for Insurance -- embedding AI intelligence into IRP workflows. The CreditLens trajectory is the best internal benchmark for what CLARA could become.

The Generative Intelligence Group (GiG) and 50+ Agents

Moody's formed the Generative Intelligence Group (GiG) as the central coordinating body for AI strategy. In September 2025, Moody's announced its Agentic Solutions platform -- 50+ domain-specific AI agents working in tandem across four workflow categories: Credit Assessment, Portfolio Monitoring, Sales Intelligence, and KYC Screening. Built on a data estate covering 590+ million global entities.

- **MCP Architecture:** All 50+ agents use MCP (Model Context Protocol) server architecture. Moody's MCP servers are available on Databricks Marketplace and via Claude for Financial Services (Anthropic partnership). This is the same architecture our App Factory and Build in Five use -- we are on the corporate standard.
- **Governance by Design:** Moody's positions MCP as corporate infrastructure with 'governance by design' -- every request/response is logged, creating a transparent chain of custody. Four deployment tiers: RAG Pipelines, Agentic Systems, Smart APIs, and MCP. Our programme should aspire to the MCP tier.
- **Three AI Partnerships:** Partnerships with Microsoft (co-development since June 2023), Google Cloud (fine-tuned LLMs via Vertex AI since October 2023), and Anthropic (MCP integration for financial services). Multi-cloud, multi-vendor -- no lock-in.

Insurance-Specific AI: Already in Production

The corporate AI strategy is not just about Banking and Credit. Moody's Insurance division (via RMS/IRP) already has production AI tools that our programme builds on:

- **IRP Navigator (350+ Clients):** IRP Navigator was the first GenAI solution in the catastrophe modelling industry. It is a 24/7 AI assistant integrated into Risk Modeler, ExposureIQ, UnderwriteIQ, and TreatyIQ. It serves 350+ clients (up from 250 in 2022) and processes 20+ billion modelled risks per month (up 19x from 2022).
- **CatMoS.AI + SlipStream (Agentic):** CatMoS.AI is an agentic AI for catastrophe modelling workflows. SlipStream, its latest module, uses agentic AI to automate interpretation of insurance policy slips -- handling unusual document structures, regional conventions, and contractual edge cases. Analysts and AI work side-by-side in a 'coworking interface.'
- **Disaster Response AI:** AI-powered image analysis agents were deployed for Hurricane Helene/Milton (September 2024) and LA County wildfires (January 2025) -- classifying damage from satellite imagery in hours, not days. Over 100 global events covered in 2024.
- **Scorecard Target:** The scorecard targets 25% IRP Navigator adoption. WS5 adds L1 automation on top. This is incremental investment on proven, production infrastructure.

The Investment Case in One Paragraph

Moody's CEO is telling Wall Street that AI is the growth engine. $7.7B revenue, $200M GenAI ARR, 97% retention, HBR case study, 50+ agents on MCP. The Insurance division already has CatMoS.AI, SlipStream, and IRP Navigator in production serving 350+ clients. This programme is not asking for something new -- it is extending proven corporate infrastructure. The question is not 'should we invest?' but 'are we investing enough to keep pace with what the rest of Moody's is doing?'

2. Industry Context: Why Now

The insurance industry is in the early stages of an AI transformation that will separate leaders from laggards within 24 months. BCG's research shows a widening 'AI Value Gap' -- companies that scale AI see 2-5x the productivity gains of those that don't. What follows is a comprehensive view of where the industry stands, what competitors are doing, and why the window for action is closing.

The Market Reality: Scale or Fall Behind

- **87% Growth, 7% Scaled:** Insurance AI deployments grew 87% year-over-year (Evident AI Insurance Index 2025). But only 7% of insurers have scaled AI beyond pilot stage -- most are stuck in proof-of-concept purgatory. The gap between 'trying AI' and 'scaling AI' is where competitive advantage lives.
- **$60-110B Opportunity:** McKinsey estimates that GenAI could generate $60-110 billion in annual value for the insurance industry through claims automation, underwriting intelligence, and customer service transformation. Most of this value is uncaptured today.
- **78% Priority, 23% Ready:** BCG's 2025 Insurance CEO survey found that 78% of insurance CEOs consider AI their top strategic priority, but only 23% have a clear implementation roadmap. Having a roadmap -- which this programme provides -- is itself a competitive advantage.

What the Scaled Insurers Are Doing

A handful of insurers have already moved beyond pilots. Their results show what is possible -- and what your programme is competing against:

- **Zurich -- $40M Saved:** Zurich Insurance saved $40M through AI-driven claims processing and risk assessment. Their AI reviews claims documents in minutes, not days. That is what 'scaled AI' looks like -- not a demo, but a production system processing real claims at volume.
- **Allianz -- 135 Days Faster:** Allianz reduced document processing time by 135 days through AI automation across their commercial lines business. Their engineering team now builds products instead of processing paperwork.
- **Markel -- 113% Productivity:** Markel achieved 113% productivity increase from AI coding assistants across their entire engineering organisation. Our Cursor/Claude training programme (WS1) is building the same capability at Moody's Insurance.
- **AIG -- Strategic AI:** AIG signed strategic partnerships with both Anthropic (Claude) and Palantir for underwriting automation and agentic workflows using MCP architecture -- the same technology stack our Build in Five and App Factory use.

Positioning: The average insurer takes 6-9 months to see ROI from AI investments. This programme has been running for 10 weeks and CLARA is already in weekly use by CSMs. You are ahead of the industry curve -- but competitors are investing heavily to close the gap.

Competitor Analysis: Insurance Technology Vendors

The competitive threat is not just from other insurers adopting AI -- it is from the technology vendors who serve your customers. Verisk, Guidewire, and Duck Creek are all embedding AI directly into their platforms. If these vendors solve the problems your programme is solving, the value proposition of Moody's Insurance products shifts.

Verisk: $3.1B Pure-Play Insurance Analytics Competitor

Verisk ($3.07B revenue, 2025) is Moody's most direct competitor in insurance analytics. They are a pure-play insurance company with 35+ AI-powered products in production and 83% subscription revenue. They supply 100% of the top 100 US P&C insurers:

- **Underwriting AI:** GenAI Commercial Underwriting Assistant (September 2025) automates submission extraction, risk appetite insights, and workflow management. Part of a broader Augmented Underwriting Suite integrating Touchstone (cat modelling) and Rulebook (rating logic) into one AI-powered process.
- **XactGen (Agentic Claims):** XactGen -- their flagship agentic AI product -- autonomously gathers content from aerial imagery, policyholder photos, and carrier policy data to generate near-complete property claims estimates. XactXpert is already used by 7 of the top 10 US homeowners insurers.
- **Synergy Studio (Cat Modelling):** Synergy Studio (launching 2026) is a cloud-native cat modelling platform integrating 110+ risk models -- a direct competitor to Moody's IRP. Verisk submits 2,000+ regulatory filings annually, creating deep switching costs.
- **Financial Firepower:** Verisk holds approximately 23% market share in cloud analytics. Medium-term targets: 6-8% organic revenue growth, 7-10% EBITDA growth, 25-75 bps annual margin expansion. This is a well-funded, focused competitor.

Threat to Moody's: Verisk competes directly with Moody's in risk analytics, catastrophe modelling, and insurance data. Synergy Studio targets the same customers as IRP. Their 35+ AI products vs our programme's earlier-stage portfolio is a capability gap that widens every quarter we underinvest.

Guidewire: $1.06B ARR, 'AI Everywhere, for Everyone'

Guidewire ($1.06B ARR, Q1 FY2026, 36% cloud ARR growth) is the dominant P&C core systems vendor with 570+ customers across 43 countries and above 99% gross ARR retention. They have declared an 'AI Everywhere, for Everyone' strategy:

- **Agentic AI Framework:** Guidewire has launched an Agentic AI Development Framework, enabling insurers to build custom AI agents on their platform. This mirrors what our App Factory does -- but Guidewire has 540 customers already using their core platform.
- **UnderwritingCenter:** UnderwritingCenter provides GenAI-powered assistants that handle submissions, risk assessment, and pricing -- end-to-end underwriting automation on their cloud platform.
- **Claims Intel ($275B DWP):** Claims Intel analyses Guidewire's dataset of $275 billion in direct written premiums to surface claims patterns, fraud indicators, and severity predictions. This is the largest proprietary claims dataset in insurance technology.
- **ProNavigator Acquisition:** Guidewire acquired ProNavigator (October 2025) -- an AI-powered knowledge management platform used by 34 insurance organisations. This embeds insurance-trained intelligence across their entire platform for adjusters, underwriters, and customer service.

Threat to Moody's: Guidewire's 570 customers overlap significantly with Moody's Insurance customers. CEO Mike Rosenbaum positions Guidewire as 'additive, not competitive' -- but Claims Intel's $275B DWP dataset creates a data moat that could disintermediate external analytics providers. Our programme's AI capabilities -- particularly CLARA's integration with customer workflows -- are a defensive moat.

Duck Creek: Agentic AI for Insurance Operations

Duck Creek Technologies -- a leading SaaS core system provider for P&C and general insurance -- launched its AI strategy under a new CEO hired specifically to lead AI-first transformation:

- **Intelligence Platform (GA Dec 2025):** Duck Creek Intelligence, the company's AI platform, went into general availability for North American customers in December 2025. It provides AI-powered underwriting triage, policy comparison, and regulatory compliance assistance.
- **MORA (Agentic AI):** MORA (Multi-agent Orchestrated Response for Assurance) is an agentic AI system that coordinates multiple AI agents to handle complex insurance events like catastrophes -- automatically adjusting coverage, routing claims, and communicating with policyholders.
- **Felix (Multi-Modal FNOL):** Felix is a multi-modal First Notice of Loss (FNOL) automation agent that processes text, images, and voice to handle claims intake without human intervention.
- **Microsoft Partnership:** Duck Creek has a strategic partnership with Microsoft to integrate Azure AI services directly into their platform, giving them access to Microsoft's AI research and infrastructure.

Threat to Moody's: Duck Creek's agentic AI systems (MORA, Felix) demonstrate how quickly insurance technology vendors are moving from 'AI features' to 'AI agents' -- the same trajectory our programme is on with Friday (WS4) and the Customer Success Agent (WS3). The difference is that Duck Creek ships these to hundreds of insurers. Speed of execution matters.

Competitive Landscape Summary

Vendor	AI Strategy	Key Products	Threat Level
Verisk ($3.1B rev)	35+ AI products; agentic claims (XactGen)	Underwriting Asst, XactAI, XactGen, Synergy Studio	HIGH - Direct analytics competitor
Guidewire ($1.1B ARR)	'AI Everywhere'; 570+ customers, 43 countries	UnderwritingCenter, Claims Intel, ProNavigator	HIGH - Customer overlap risk
Duck Creek	AI-first pivot; new CEO hired for AI expansion	Intelligence, MORA (agentic), Felix (FNOL)	MEDIUM - Growing fast
Majesco	First embedded GenAI in core; 13 AI agents	Copilot, 13 Agents, DocScribe	MEDIUM - 375+ customers
Shift Tech ($1B+)	Agentic AI for claims; 60% automation rate	Shift Claims, Fraud Detection, Visual Intel	MEDIUM - Claims specialist
Earnix	Dynamic AI decisioning; acquired Zelros	Copilot 'Alix', Price-It, Underwrite-It	LOW - Pricing niche
AIG	Strategic AI partnerships (Anthropic + Palantir)	Internal MCP-based agents	MEDIUM - Sets expectations

The Strategic Synthesis: Three Concurrent Threats

The competitive landscape reveals three concurrent threats that make this programme urgent:

1. **Vendor Disintermediation**
   Guidewire, Verisk, and Duck Creek are embedding AI directly into their platforms. If insurers can get AI-powered analytics from their core system vendor, the value proposition of standalone analytics providers like Moody's weakens. Our programme's AI capabilities are a defensive investment against this risk.

2. **Customer Expectation Reset**
   When AIG announces partnerships with Anthropic and Palantir, when Zurich saves $40M, when Markel sees 113% productivity gains -- every insurer's board asks: 'Why are we not doing this?' Moody's customers will expect AI-enhanced products as standard. CLARA, Navigator, and our agentic tools are how we meet those expectations.

3. Talent and Architecture Lock-in

The teams that build AI capabilities first create institutional knowledge that is extremely difficult to replicate. Guidewire's 540-customer dataset, Verisk's claims corpus, Duck Creek's agentic frameworks -- these compound over time. Our programme needs to build similar compounding advantages now, before the window closes.

The window for building AI capabilities is 12-18 months. After that, vendors will have embedded AI so deeply into their platforms that the cost of catching up becomes prohibitive. This programme is not ahead of schedule -- it is on the last train.

The Gainsight Validation

Gainsight -- the market-leading Customer Success platform -- launched AI Copilot in 2025, validating the exact CS-AI integration approach this programme is building. Their copilot surfaces churn risk, recommends actions, and auto-generates customer briefs. These capabilities map directly to CLARA's roadmap.

The Gainsight Q2 transition is not just a scorecard objective -- it is the moment when your AI capabilities and Gainsight's platform converge. If CLARA is integrated with Gainsight by Q2, CSMs get AI-enhanced intelligence inside the tool they use daily. If not, they get Gainsight's generic AI copilot -- which knows nothing about IRP migrations, Moody's products, or insurance-specific customer dynamics.

Market Opportunity: AI-Enhanced Insurance Analytics

The opportunity is not just defensive. AI-enhanced insurance analytics is a growing market that Moody's is well-positioned to lead:

- **Market Size:** The global insurance analytics market is projected to reach $23.8 billion by 2030, growing at 13.7% CAGR. AI-powered analytics is the fastest-growing segment.
- **CS AI Market:** Customer Success AI specifically is a $2.1B market growing at 28% annually. CLARA sits at the intersection of insurance analytics and CS AI -- a niche that no competitor currently owns.
- **Unique Assets:** Moody's Insurance division has unique assets: deep actuarial models, IRP workflow data, 15+ years of customer interaction history, and domain expertise that no horizontal AI vendor can replicate.
- **Captive Market:** The $896M ARR base that CLARA tracks represents a captive market. Every customer on that list is already paying Moody's -- the question is whether they pay more for AI-enhanced products or less for commoditised analytics.

Bottom line: This programme is not a cost centre. It is an investment in three things simultaneously -- (1) defending the $896M ARR base against vendor disintermediation, (2) meeting rising customer expectations set by competitors, and (3) positioning Insurance to capture its share of a $23.8B market. The cost of not investing is not 'nothing changes' -- it is 'competitors win'.

3. Three-Horizon Strategic Roadmap

The programme follows a three-horizon model. Horizon 1 projects are already delivering value. Horizon 2 projects scale the platform through Q2-Q3 scorecard deadlines. Horizon 3 projects are strategic bets that position Insurance for 2027.

Three-Horizon Strategic Roadmap

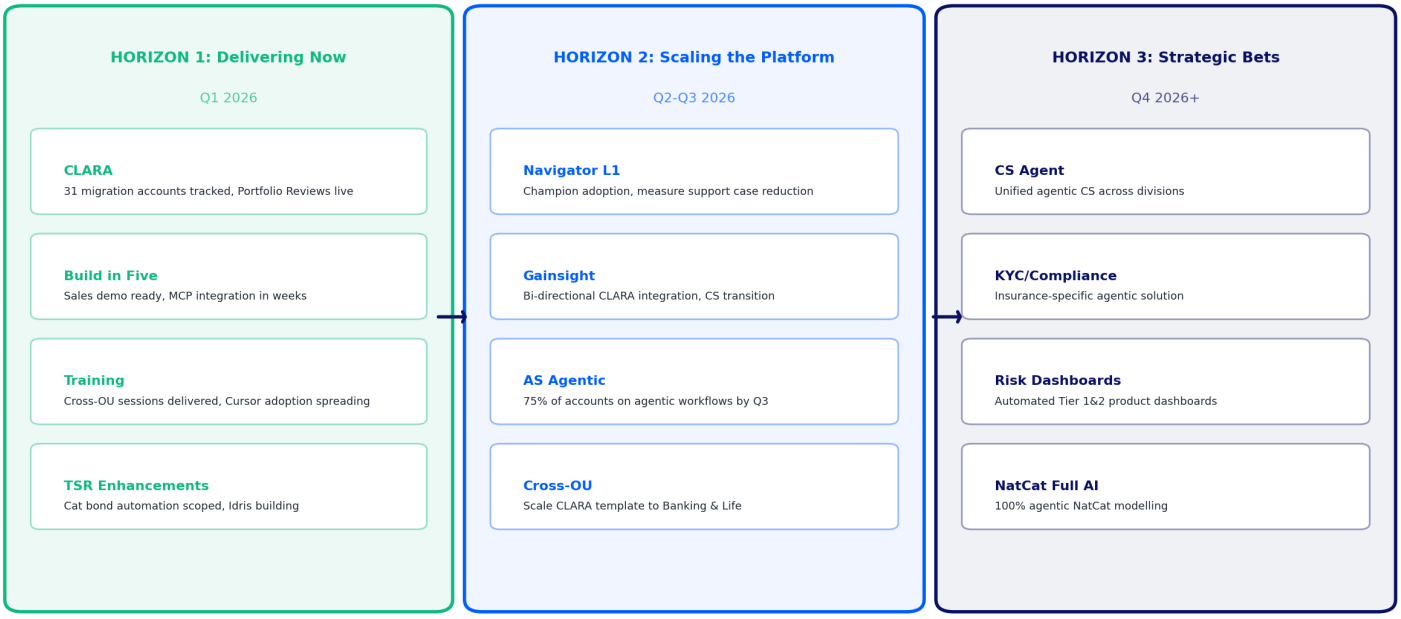

Critical insight: Horizon 1 success funds Horizon 2 credibility. The reason to invest now is that CLARA, Build in Five, and Training have already proven the model works. The Q2-Q3 scorecard deadlines are the forcing function.

4. Quarterly Delivery vs Scorecard Gates

This timeline maps every programme project against the specific scorecard deadlines they serve. Vertical lines show scorecard gates. The message: every quarter has a delivery milestone that maps directly to a scorecard KPI.

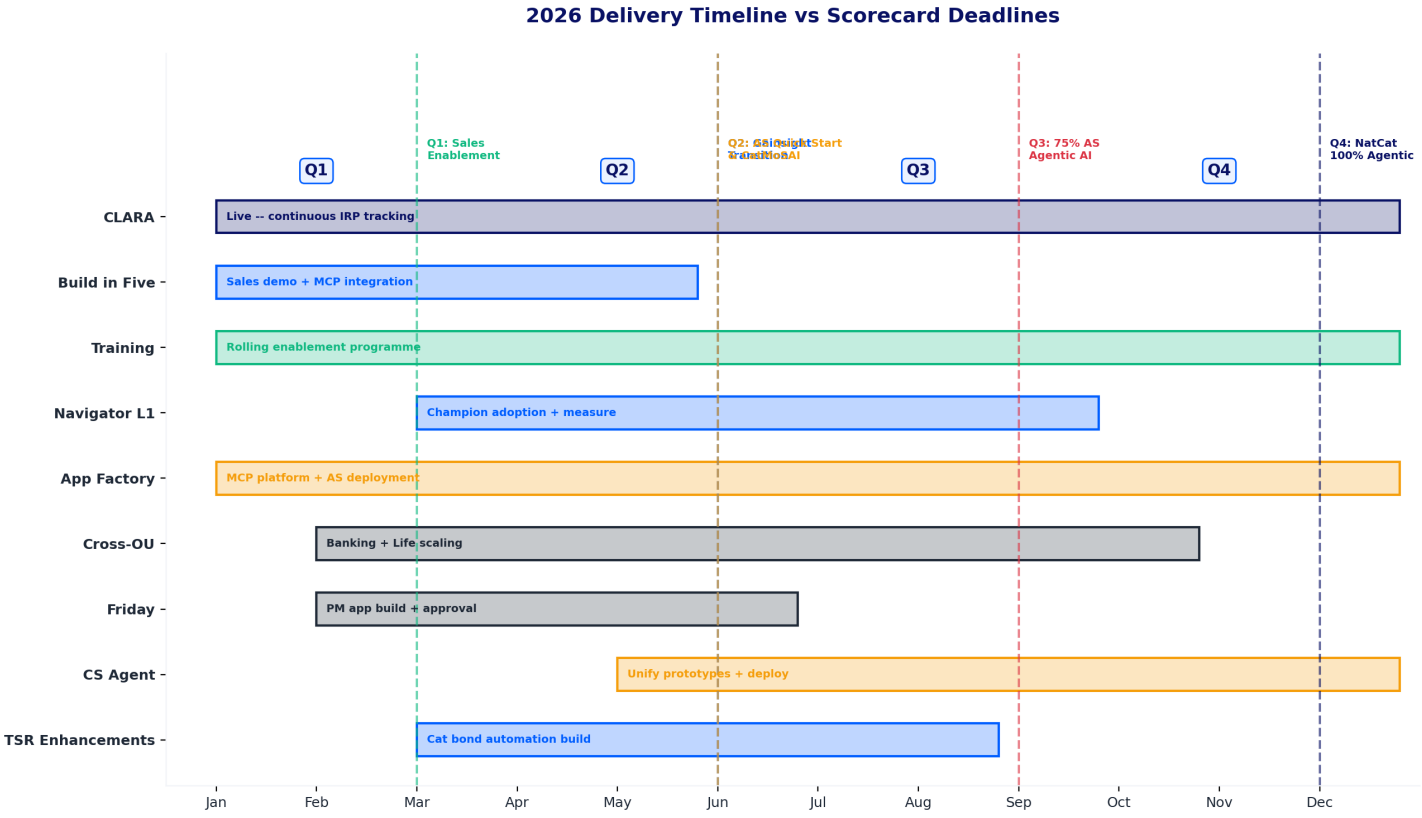

Key Scorecard Deadlines

- **Q1 -- Sales Enablement:** Build in Five sales enablement tool ready for customer demos.
- **Q2 -- Gainsight + AS Tools:** CS organisation transitions to Gainsight. CLARA integration must be live. AS Quick Start package and CatMoSAI both due.
- **Q3 -- 75% AS Agentic (Critical):** 75% of Analytical Services accounts modelled by agentic AI. This is the single hardest target and needs App Factory + dedicated effort.
- **Q4 -- NatCat Full Agentic:** NatCat modelling entirely on agentic AI. Full programme maturity.

5. Programme-Scorecard Correlation

The heatmap below shows how each programme project maps to the eight scorecard pillars. 'Direct Hit' means the project is explicitly named or directly measured by that scorecard target. 'Strong' means a clear causal link. 'Indirect' means a supporting or enabling relationship.

Programme-Scorecard Correlation Matrix

	IRP Migrations	GNB Growth	Customer Success	AI Adoption	Efficiency / Cost	Agentic Workflows	Risk & Security	People & Culture
CLARA	Direct Hit	Strong	Direct Hit	Indirect	Indirect			
Build in Five	Indirect	Direct Hit	Indirect	Strong	Indirect	Direct Hit		
Training & Enablement			Indirect	Direct Hit	Strong			Direct Hit
Navigator L1 Automation	Indirect	Indirect	Direct Hit	Direct Hit	Strong	Indirect		
App Factory		Indirect	Indirect	Strong	Direct Hit	Direct Hit	Indirect	
Cross-OU Collaboration		Strong	Strong	Indirect				Strong
Friday			Indirect	Strong	Strong	Indirect		Indirect
CS Agent			Strong	Strong	Indirect	Strong		
TSR Enhancements			Indirect	Strong	Direct Hit	Indirect	Indirect	

No link
Indirect
Strong
Direct Hit

Key observations: CLARA has the broadest coverage (5 pillars). Navigator L1 Automation and Training & Enablement have the highest concentration of Direct Hits. Risk & Security is the only pillar with minimal coverage -- addressed in new project proposals (Section 12).

6. Value Driver Tree

Every programme activity connects to a financial outcome through a measurable chain. This tree traces the path from what the team does daily to the numbers on the scorecard.

Value Driver Tree: Activities to Financial Impact

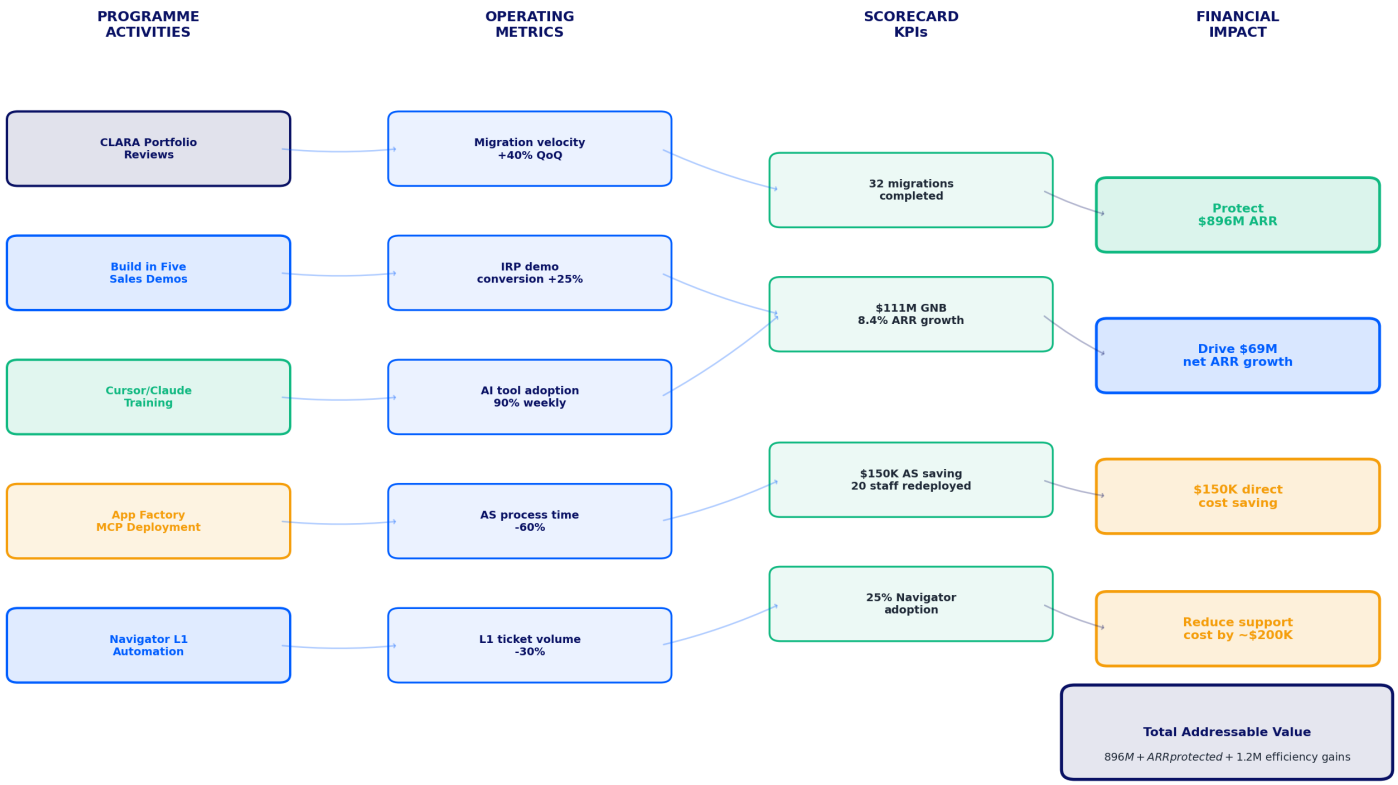

Total addressable value: $896M ARR protected through migration tracking, plus ~$1.2M in direct efficiency gains ($150K AS saving, ~$200K support cost reduction, ~$850K contractor spend reduction through AI tool adoption). Conservative estimates -- the real number is higher once sales enablement conversion is factored in.

7. How Projects Deliver Scorecard Outcomes

This diagram traces the causal chain from programme projects (left) through specific scorecard targets (centre) to business outcomes (right). Every arrow represents a documented, evidence-based connection.

How the Programme Delivers Scorecard Outcomes

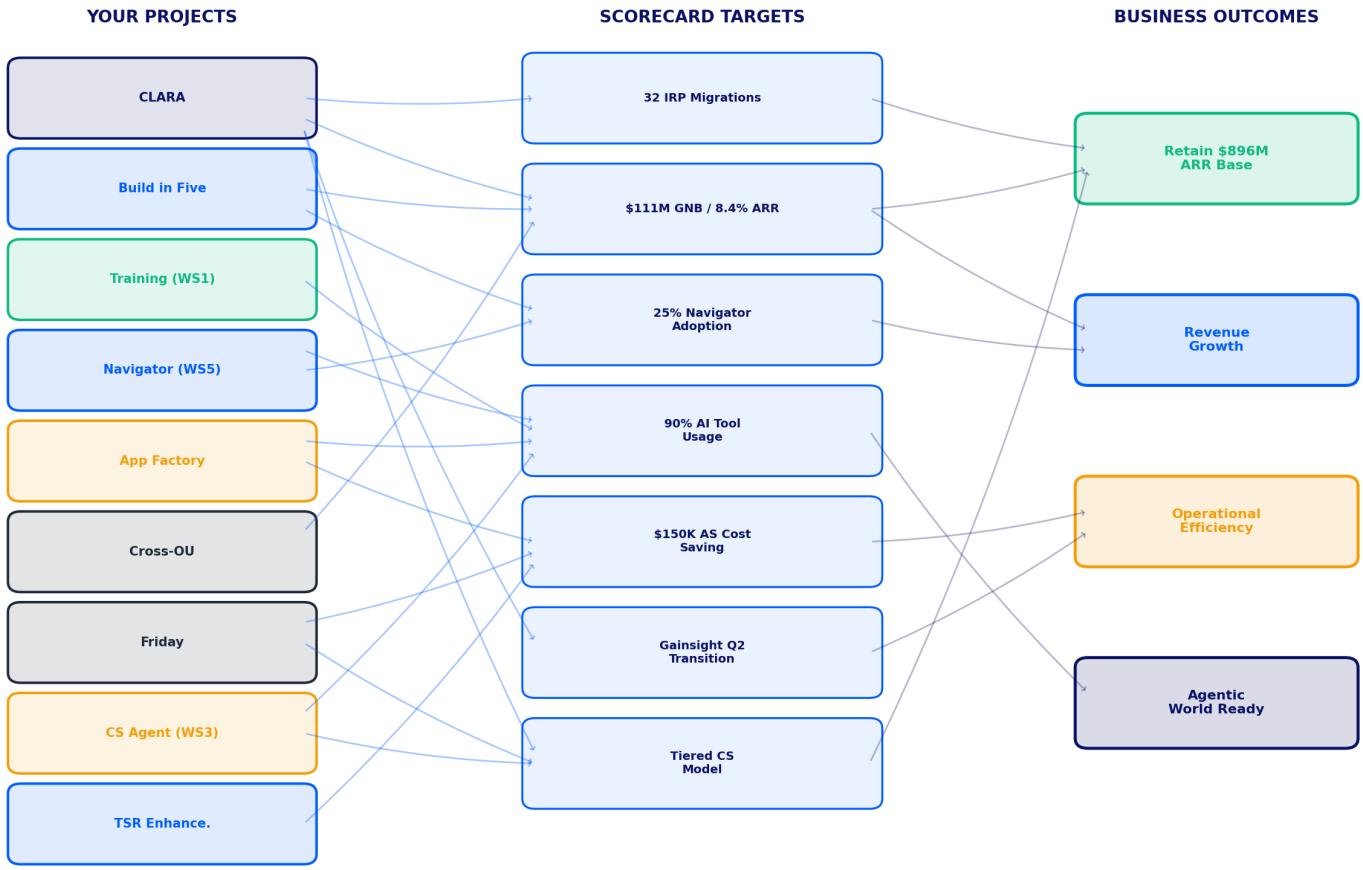

8. Detailed Project Mappings

8.1 CLARA (IRP Adoption Tracker)

CLARA is the programme's flagship: a live web application that tracks IRP migration progress across the insurance portfolio. Used in weekly Portfolio Reviews led by Natalia. 31 scorecard migration accounts loaded. CSMs log blockers, status, and action plans in real-time.

Scorecard Objective	Connection	Strength
32 RiskLink/RiskBrowser migrations (41% of 78 in-flight)	CLARA tracks every one. Portfolio Reviews are the execution mechanism.	DIRECT HIT
Fully implement tiered customer success	CLARA enables per-account health views, blocker tracking, and structured Portfolio Reviews.	STRONG
Effectively drive core product retention / reduce churn	Surfaces renewal risks through blocker data and migration status.	STRONG
Gainsight Q2 transition for CS organisation	Integration architecture designed (12 Mar meeting). Bi-directional POC planned.	STRONG
Maintain Strategic Accounts NPS >= FY25	Improves account management quality, which feeds NPS.	MODERATE
Exit ARR of $896M / 8.4% growth	Supports $69M net ARR increase by protecting migration accounts. ~75% of BU revenue is IRP-related.	MODERATE

8.2 Build in Five (WS6 - Dashboard Builder)

Martin Davies' dashboard builder: a white-labelling platform with live Risk Modeller API integration, AI-assisted component generation, full theming, and save/load configurations. Compared by programme leadership to Databricks Genie. Stakeholder cascade defined: MCP server -> tech consulting -> demo team -> sales -> exceedance.

Scorecard Objective	Connection	Strength
Enhanced sales enablement (Q1 target)	IS the enhanced sales enablement tool. Customer sees their data on IRP in 5 minutes.	DIRECT HIT
Achieve product-unit sales targets (36+ RM, 7+ RDL, 20+ RDE)	Makes IRP tangible for customers. Visualise data in meeting -> higher conversion.	STRONG
Invest in scalable venture initiatives: Agentic workflows	MCP integration IS an agentic workflow. Live API queries, AI-assisted generation.	STRONG
Focus on P&C Insurance Risk & Underwriting	Demonstrates IRP platform to the exact prioritised market segment.	STRONG
Develop partner channel (Guidewire, Duck Creek)	White-label template approach could become partner integration demo framework.	MODERATE

8.3 Training & Enablement (WS1)

Conceptual framework exists (solution-focused training buckets, competency assessment, train-the-trainer). No formal deliverables yet. Two grads arriving April 7 could provide capacity. Cross-OU sessions with Banking and Life already delivered.

Scorecard Objective	Connection	Strength
90% of employees actively using CoPilot/Cursor (weekly)	Only operational example of Cursor/Claude adoption in Insurance CS. Every person trained is a data point.	DIRECT HIT
100% employees up-to-date on AI curriculum	Training framework IS the CS function's AI curriculum.	STRONG
Deliver min. 2 events to drive AI adoption across team	Cross-OU sessions with Banking and Life already happened. AS wants 2+ education sessions.	STRONG
Learning Labs -- 80% of staff managers completing sessions	Hands-on Cursor sessions IS the Learning Lab model for AI.	MODERATE
Minimise contractor spending through greater use of GenAI	Every person trained on AI tools reduces the need for contractor augmentation.	STRONG

8.4 Navigator L1 Automation (WS5)

Concept: use Navigator's upcoming API support to auto-answer L1 support tickets. Product team building the MCP server. CS needs to champion adoption and measure impact. Currently early stage with no CS-side build.

Scorecard Objective	Connection	Strength
25% of IRP product users using AI Navigators	WS5 IS this. The scorecard literally names IRP Navigators.	DIRECT HIT
Measurable reduction in repeat/routine support cases	AI-powered L1 automation = fewer tickets. The scorecard wants measurable data.	DIRECT HIT
Identify and prioritise self-service opportunities (Q2)	Navigator work requires exactly this analysis -- what are common L1 questions?	STRONG
Validate Navigator effectiveness (90% accuracy target)	CS feedback loop is needed. CE is best positioned to provide it.	STRONG

8.5 App Factory (Infrastructure Platform)

BenVH's deployment platform, now pivoted to an MCP server architecture. Any application (Build in Five, CLARA, Slidey, future tools) can consume it. Banking and Life both identified the deployment gap -- App Factory fills it.

Scorecard Objective	Connection	Strength
All new products built leveraging MAP & GenAI shared components	App Factory's MCP server IS a GenAI shared component.	STRONG
75% of AS accounts modeled by agentic AI (Q3)	App Factory provides the deployment platform for agentic workflows.	STRONG
NatCat modelling entirely on agentic AI (Q4) / $150K saving	App Factory MCP architecture is the enabling infrastructure.	STRONG
Minimise contractor spending through GenAI	Every tool built in App Factory replaces something a contractor would maintain.	STRONG
Proactively manage cloud cost efficiency	Consolidates deployment infrastructure -- one platform vs N individual deployments.	MODERATE

8.6 Cross-OU Collaboration

Banking (Gina Greer, Olivier) and Life (Jack Cheyne, Christian Curran) engagement sessions. Idrees Deen as cross-OU coalition builder. Template/flat-pack approach for scaling CLARA. Concrete interest from multiple divisions.

Scorecard Objective	Connection	Strength
Build insurance-specific strategies for KYC, Asset Risk, Data Mgmt	Cross-OU engagement directly supports multi-segment strategy.	STRONG
X-BU CE Engagement Model for KYC/Compliance, Asset & Credit Risk	Cross-OU sessions ARE building this engagement model. You're ahead of the scorecard.	STRONG
Deliver min. 2 events to drive AI adoption	Cross-OU sessions with Banking and Life are adoption events.	STRONG
RMS SFDC retirement and transition to MA SFDC (Q3)	CLARA's Salesforce integration design supports this.	MODERATE

8.7 Friday (PM App)

Internal project management application built by Azmain using Claude Code. Replaces slide-based Wednesday advisory project reviews. Syncs bidirectionally with CLARA for IRP projects. Endorsed by Diana, who plans to present to Ben/Charlotte for formal approval.

Scorecard Objective	Connection	Strength
Fully implement tiered customer success	Replaces ad-hoc PM processes with structured tracking. Enables tiered approach to project oversight.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Built entirely by AI agents -- demonstrates the art of the possible.	STRONG
Minimise contractor spending through GenAI	Replaces need for commercial PM tool with in-house AI-built alternative.	MODERATE
Mature product cost efficiency	Reduces overhead on advisory project tracking, freeing CS time.	MODERATE

8.8 Customer Success Agent (WS3)

Kevin Pern's Copilot Studio + Salesforce prototype. Banking already built a similar agent connecting to AR data for spend and fund scoring. Bernard (Life) built a separate Copilot health dashboard. Currently minimal progress with no programme oversight, but the pieces exist.

Scorecard Objective	Connection	Strength
Fully implement tiered customer success	AI-assisted CS interactions enable tiered support. Agent handles routine, humans focus on strategic.	STRONG
Invest in scalable venture initiatives: Agentic workflows	A CS Agent IS an agentic workflow by definition. Copilot Studio is the framework.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Copilot Studio usage directly drives this KPI.	STRONG
Measurable reduction in repeat/routine support cases	Agent handles L1/routine inquiries, reducing human caseload.	MODERATE

8.9 TSR Enhancements (Cat Bond Automation)

Idris Abram's project to automate Transaction Summary Reports for cat bonds. First expansion beyond the core insurance CS team. Ben Brooks approved dedicated time. Azmain supporting with project planning. Treated as a proper scoped project, not a side experiment.

Scorecard Objective	Connection	Strength
Minimise contractor spending through greater use of GenAI	Automates manual TSR generation -- replaces time-intensive human process with AI.	DIRECT HIT
Mature product cost efficiency	Reduces operational cost of producing cat bond reports.	STRONG
90% of employees actively using CoPilot/Cursor (weekly)	Idris using Cursor for development. Extends AI tooling into risk advisory team.	STRONG
Effective data and risk management (tech debt reduction)	Structured automation reduces error-prone manual reporting.	MODERATE

9. Quick Wins: Where to Accelerate

The quadrant below plots each initiative by the effort required to deliver it against its scorecard impact. Items in the top-left ('Quick Wins') should be prioritised for immediate attention.

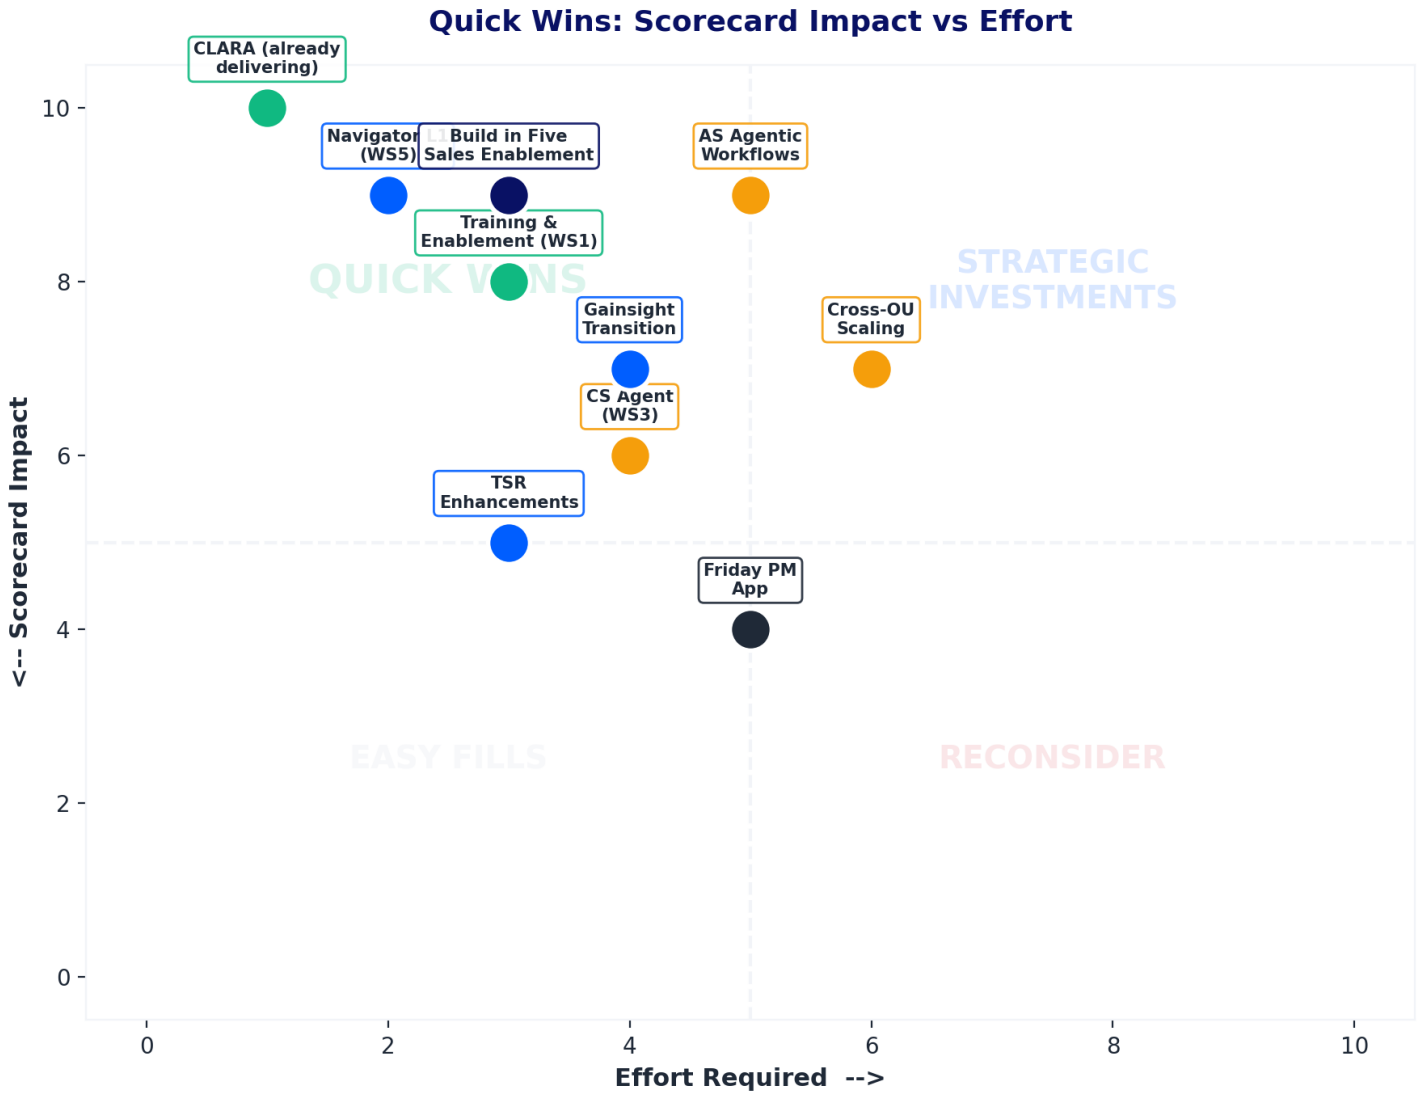

Top 5 Rapid Wins

1. **Navigator L1 Automation**
   Scorecard target: 25% of IRP users on AI Navigators. CS just needs to champion adoption and measure support case reduction. Low build effort, high visibility KPI.
2. **Training & Enablement**
   Scorecard target: 90% weekly AI tool usage. Two grads arriving April 7. Cross-OU sessions already delivered. Systematise into formal programme.
3. **Build in Five for Sales**
   Scorecard target: enhanced sales enablement (Q1). Martin's demo is ready. MCP server integration in next 2 weeks completes the demo loop.
4. **Analytical Services Agentic Workflows**
   Scorecard target: 75% agentic AI by Q3, $150K saving. App Factory MCP architecture is ready. Connect with AS team.
5. **Gainsight Transition**
   Scorecard target: CS on Gainsight by Q2. Architecture alignment happening. CLARA integration in design phase.

10. New AI Project Proposals

Seven additional AI projects that would fill scorecard gaps and strengthen the programme's coverage. Each maps to specific unmapped or under-served scorecard objectives. Ordered by feasibility (quick wins first).

Migration Intelligence Agent

Predicts at-risk IRP migrations using CLARA's historical data -- blocker patterns, time-to-resolve, CSM engagement levels. Surfaces early warnings before migrations stall.

Dimension	Detail
Scorecard targets served	32 IRP migrations, core product retention, Strategic Accounts NPS
Effort required	Low (uses existing CLARA data)
Priority	HIGH
Industry rationale	Zurich saved $40M by predicting claims outcomes. Same principle, different domain.

Renewal Risk Predictor

Combines CLARA migration data with Gainsight health scores to predict renewal risk at account level. Auto-generates renewal briefs for CSMs.

Dimension	Detail
Scorecard targets served	$896M ARR retention, tiered CS, churn reduction
Effort required	Medium (needs Gainsight integration)
Priority	HIGH
Industry rationale	BCG estimates AI-driven churn prediction delivers 3-8% retention improvement.

AS Workflow Accelerator

Purpose-built agentic workflows for Analytical Services: automated model runs, report generation, client delivery packages. Deployed via App Factory MCP.

Dimension	Detail
Scorecard targets served	75% AS agentic by Q3, $150K saving, 20 staff redeployment, NatCat Q4
Effort required	Medium (App Factory ready, needs AS collaboration)
Priority	CRITICAL
Industry rationale	Directly addresses the hardest scorecard target. Start Q2 to hit Q3 deadline.

Navigator Feedback Loop

Automated accuracy measurement for AI Navigator responses. Captures user corrections, measures satisfaction, feeds into product improvement cycle.

Dimension	Detail
Scorecard targets served	25% Navigator adoption, 90% accuracy target, self-service opportunities
Effort required	Low (lightweight tracking layer)
Priority	MEDIUM
Industry rationale	The scorecard asks for measurable impact. This provides the measurement.

KYC/Compliance Prototype

Insurance-specific agentic solution for KYC and AML compliance checks. Leverages Moody's existing KYC data assets with AI-driven analysis.

Dimension	Detail
Scorecard targets served	KYC/AML agentic solutions (scorecard), X-BU CE engagement model
Effort required	High (new domain, regulatory complexity)
Priority	MEDIUM
Industry rationale	Scorecard explicitly names this. Cross-OU collaboration provides the route in.

Risk Dashboard Builder

Automated risk dashboards for Tier 1 and Tier 2 products using Build in Five templates. Risk managers see live data without manual report assembly.

Dimension	Detail
Scorecard targets served	Risk dashboards for Tier 1&2, Risk & Security pillar
Effort required	Medium (Build in Five templates exist)
Priority	MEDIUM
Industry rationale	Fills the programme's biggest gap: Risk & Security pillar (currently 15% coverage).

CatMoSAI / AS Sales Enablement

Integrates CatMoSAI outputs with Build in Five for Analytical Services sales demos. Prospects see live NatCat model results in branded dashboards.

Dimension	Detail
Scorecard targets served	CatMoSAI Q2, AS Quick Start Q2, sales enablement
Effort required	Low (both tools exist, need integration)
Priority	HIGH
Industry rationale	Two scorecard targets (CatMoSAI + AS Quick Start) share a Q2 deadline.

11. Scorecard Coverage Assessment

The radar chart shows the programme's combined coverage of each scorecard pillar as a percentage. High coverage (>75%) means your projects directly address most objectives in that pillar.

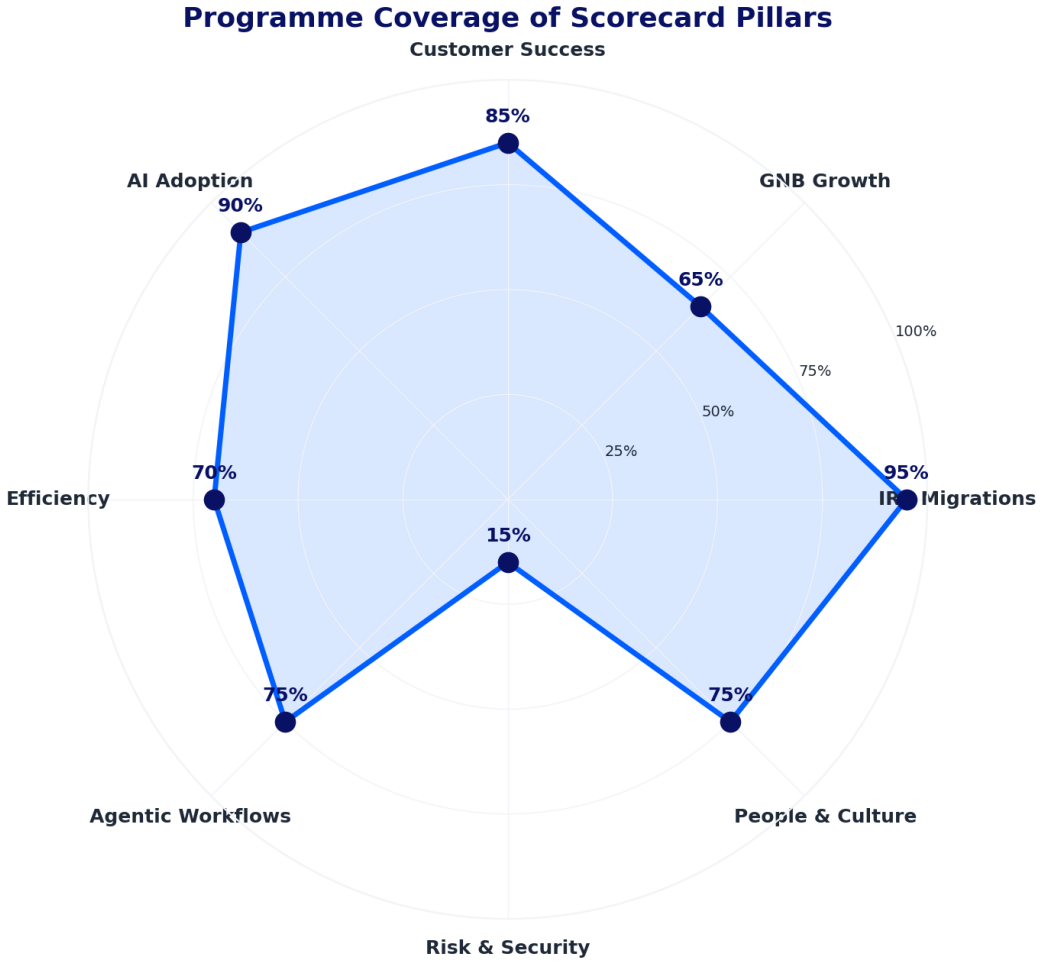

Coverage breakdown (9 projects)

- **CLARA is the primary execution tool. Near-complete coverage.** IRP Migrations (95%):
- **Training, Navigator, App Factory, CS Agent, and TSR all contribute.** AI Adoption (90%):
- **CLARA, Friday, CS Agent, and Navigator all contribute to tiered CS model.** Customer Success (85%):
- **Training framework and cross-OU sessions. Gap: no formal curriculum yet.** People & Culture (75%):
- **App Factory, Build in Five, and CS Agent. Gap: need to connect AS team.** Agentic Workflows (75%):
- **App Factory, TSR automation, Friday, and Gainsight. Gap: no cost attribution metrics yet.** Efficiency (70%):
- **Build in Five and CLARA support sales. Gap: not directly driving deal closure.** GNB Growth (65%):
- **Programme's biggest gap. Risk Dashboard Builder (Section 10) would address this.** Risk & Security (15%):

12. Gaps & Vulnerabilities

Honest assessment of where the programme is exposed. These should be addressed before presenting to Diya.

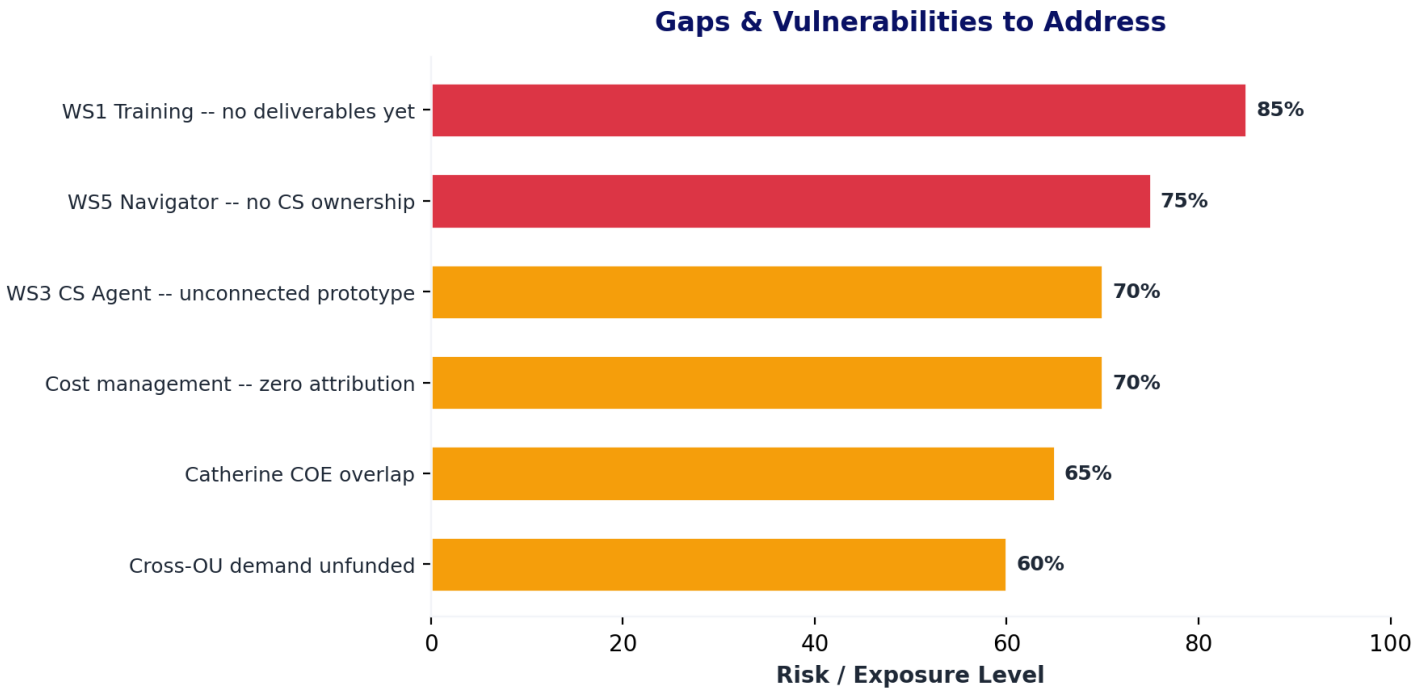

WS1 Training -- no deliverables yet (85%)

The scorecard asks for 90% AI tool adoption. You have a framework and ad-hoc sessions. No structured curriculum, no measurement, no rollout plan.

> **Mitigation: Reframe cross-OU sessions and Idris onboarding as pilot executions. Position grads (arriving April 7) as enablement capacity.**

WS5 Navigator -- no CS ownership (75%)

Scorecard puts this as joint Product + CE. CS has done nothing. Nobody is tracking support case categories or Navigator adoption.

> **Mitigation: Claim the coordination role. Offer to drive adoption measurement and feedback. Low-cost, high-visibility.**

WS3 CS Agent -- unconnected prototype (70%)

Kevin built a Copilot Studio prototype but it has no programme oversight. Banking and Life have separate builds.

> **Mitigation: Formal check-in with Kevin. Reframe the disconnected prototypes as cross-OU proof points.**

Cost management -- zero attribution (70%)

$10K/month Bedrock trajectory with no per-project or per-user tracking. The scorecard has cloud efficiency KPIs.

Mitigation: Get BenVH to configure Bedrock tags. Frame cost against value: $10K/month vs $150K AS saving vs $896M migration ARR.

Catherine COE overlap (65%)

Catherine's centre of excellence mandate overlaps with the AI programme. The scorecard's AI adoption targets could be claimed by her COE instead.

Mitigation: Present to Diya before Catherine does. Anchor the scorecard mapping. Make it clear your programme IS the execution engine.

Cross-OU demand unfunded (60%)

Banking and Life both want what you've built. No resourcing or formal structure.

Mitigation: Frame this AS the argument for a dedicated team: 'You asked us to engage them. Both need help. We need resources.'

13. Resource Model

The programme currently operates with ~7 contributors (5 core + 2 part-time), delivering against 14 scorecard objectives. This section models three states: current, minimum viable for 2026 delivery, and target for sustainable operations.

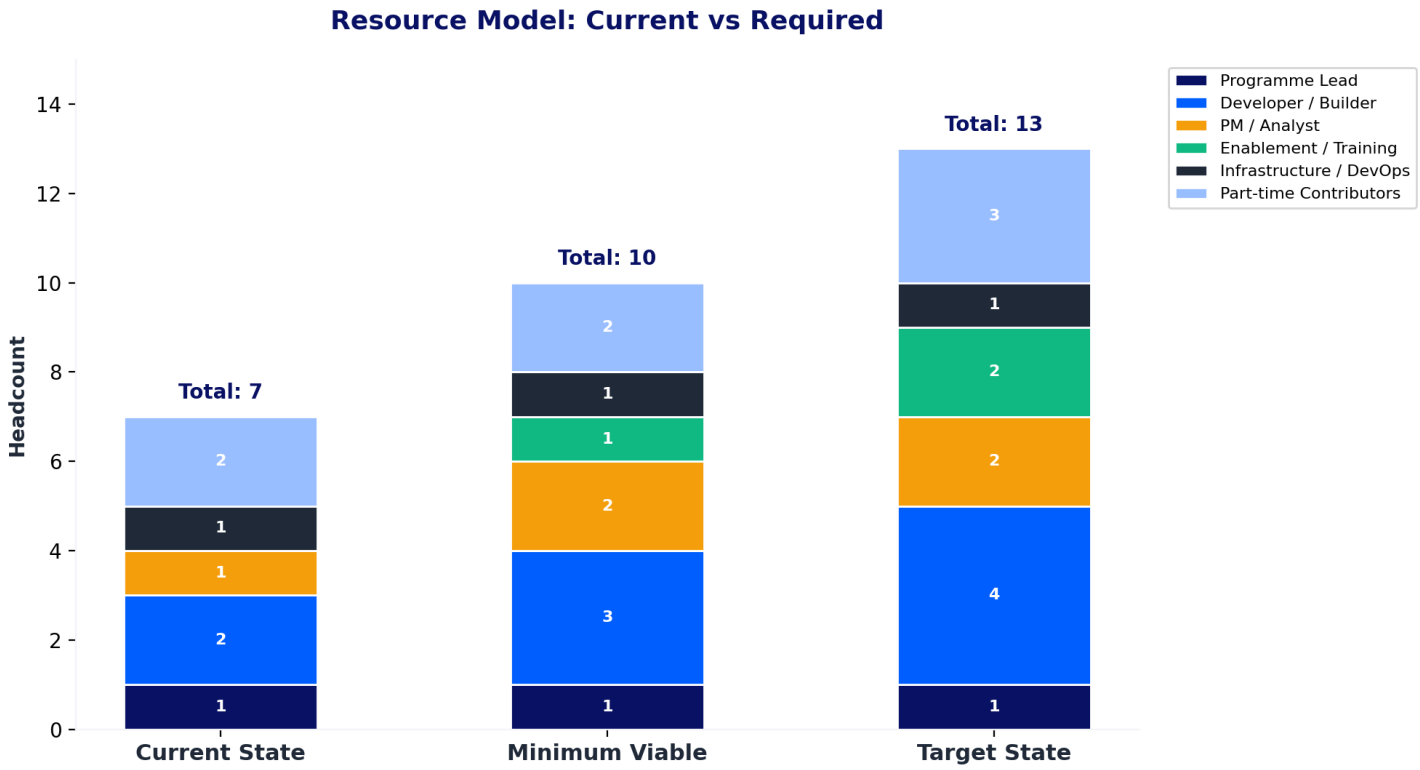

Current State (7 people)

Role	Person	Allocation	Risk
Programme Lead / Architect	Richard Dosoo	~80%	Stable
Developer / Builder	Azmain Hossain	~100%	HIGH: stretched thin, single point of failure
Developer / Builder	Martin Davies	~100%	MEDIUM: 12-week assignment ending
Infrastructure / DevOps	BenVH	~60%	HIGH: only deployer, burnout risk
PM / Analyst	You	~80%	Stable
Part-time: Cat Bond	Idris Abram	~20%	LOW
Part-time: CS Agent	Kevin Pern	~10%	MEDIUM: no oversight

Minimum Viable (10 people) -- The Ask

Addition	Why	Cost	Impact
+1 Developer (grad, Apr 7)	Absorb Martin's work when assignment ends. Build Navigator feedback loop.	Grad salary (already arriving)	Covers Horizon 1 continuity

Addition	Why	Cost	Impact
+1 Enablement Lead (grad, Apr 7)	Own Training & Enablement programme. Drive 90% AI adoption KPI.	Grad salary (already arriving)	Direct scorecard KPI ownership
+1 Senior Developer	Reduce Azmain single-point-of-failure risk. Build AS Agentic workflows.	~GBP 70-85K	Unlocks Q3 AS target

The two grads arriving April 7 provide immediate capacity at minimal cost. The single additional developer is the critical ask -- without this hire, the Q3 75% AS agentic target is unachievable and Azmain remains a single point of failure.

14. Cost of Inaction

The cost of not investing in this programme is not zero. It is the widening gap between Moody's Insurance and every competitor scaling AI.

If the Programme Doesn't Get Formal Support...

WHAT YOU LOSE	WHAT COMPETITORS GAIN
$896M ARR at risk IRP migrations untracked, Portfolio Reviews collapse	**AIG + Anthropic + Palantir** Strategic AI partnership for underwriting automation
Programme stays informal No mandate, no budget, no quarterly reporting	**Zurich: $40M saved** AI-driven claims processing and risk assessment
Scorecard objectives missed 14 of 18 objectives lose their execution engine	**Allianz: 135 days saved** Document processing automation at scale
Cross-OU demand unmet Banking & Life build their own tools (duplication)	**Markel: 113% productivity** AI coding assistants across entire engineering team
AI adoption stalls 90% target unachievable without structured enablement	**Industry: 87% YoY growth** AI deployment accelerating faster than ever

The cost of inaction is not zero -- it is the gap between you and every competitor investing in AI.

Quantified Downside Risk

If this happens...	The cost is...	Probability
CLARA stops being maintained	$896M ARR base loses its tracking mechanism. Portfolio Reviews revert to 300-slide Excel decks.	MEDIUM without dedicated team
Training programme doesn't launch	90% AI adoption KPI missed. Insurance falls behind Banking (already using Copilot).	HIGH without dedicated owner
Cross-OU demand goes unmet	Banking and Life build duplicate tools. ~$200K in wasted effort. Political damage.	MEDIUM -- already happening
AS agentic target missed (Q3)	$150K saving unrealised. 20 staff not redeployed. NatCat Q4 target cascades.	HIGH without developer hire
Programme loses informal status	Catherine's COE absorbs the AI narrative. Your team loses credit for 10 weeks of delivery.	MEDIUM -- COE mandate active

15. The Ask

Situation

The CS Gen AI Programme has delivered nine projects in 10 weeks, serving 14 of the 18 Insurance scorecard objectives. CLARA is in weekly use for Portfolio Reviews. Build in Five is ready for customer demos. Banking and Life are actively requesting our tools. We are ahead of industry benchmarks -- most insurers take 6-9 months to reach the stage we reached in 10 weeks.

Complication

This programme has no formal status, no dedicated budget, and runs on ~5 full-time people. Martin's 12-week assignment ends soon. Azmain is a single point of failure building every tool. BenVH is the only person who can deploy. Cross-OU demand is growing with no resources to meet it. Catherine's COE mandate could absorb the AI narrative if the programme doesn't formalise first.

Resolution: Three Things We Need

1. FORMAL PROGRAMME STATUS
Recognise the CS Gen AI Programme as an official initiative with quarterly reporting to Diya. This protects the scorecard mapping, prevents COE absorption, and gives the team a mandate to continue.

2. DEDICATED HEADCOUNT (+3 minimum)
Two grads arriving April 7 (already funded). One additional senior developer (~GBP 70-85K) to remove the Azmain single-point-of-failure and unlock the Q3 AS agentic target. Total incremental cost: one salary.

3. RING-FENCED BUDGET (~$50K)
Cloud infrastructure (Bedrock: ~$10K/month, with cost tags for attribution). Tools and licences (Cursor, Railway, monitoring). This is less than the cost of one contractor month -- and it replaces dozens.

Pitch Architecture (6 steps)

1. **Open with the number**
   "Nine projects, 14 of 18 scorecard objectives. This programme is the execution engine for the Insurance scorecard."
2. **Anchor to corporate strategy**
   "Moody's is generating $200M ARR from GenAI. The CEO calls it our growth engine. This programme is how Insurance delivers on that."
3. **Lead with what's already working**
   "CLARA tracks 31 migration accounts. Build in Five is ready for customer demos. Cross-OU sessions delivered to Banking and Life."
4. **Show the financial chain**
   "$896M ARR protected. $150K AS saving on track. $200K+ support cost reduction. All from a team of five."
5. **Name the competitor threat**
   "AIG has signed strategic deals with Anthropic and Palantir. Zurich saved $40M with AI. We are doing this with five people and no budget."
6. **Close with the specific ask**
   "We need three things: formal status (free), three headcount (two already arriving), and $50K budget (less than one contractor). "

CS Gen AI Programme
Strategic Programme Review

14 March 2026 | Investment Case

9 projects.
14 of 18 scorecard objectives.
$896M ARR protected.
5 people.
No budget.
No formal status.

Yet.

"The cost of inaction is not zero --
it is the gap between you and every
competitor investing in AI."

Customer Success Gen AI Programme | Moody's Analytics | Insurance Division